

Process Address Space

What is operating system?

- ❑ Operating system is a software that provides an easy to use interface between user and machine.
 - ❑ Example: Copying one file to another
 - ❑ Using OS requires a simple command or right click and select “copy and paste” option
 - ❑ Without OS: an user needs to
 - Read data from hard disk blocks (first file A) into memory
 - Write data to other disk blocks for new file B
 - Needs to update directory structure so that blocks related to other files are not overwritten
-

OS Functions

- ❑ There are three main functions that an OS performs
 - Provides an easy-to-use interface between user and machine. User need not understand functioning of hardware. Irrespective of underlying hardware, interface to user remains the same.
 - Manages computer resources on behalf of user. Resources are hardware (CPU/Memory/IO device/Netowrk interface cards/GPGPU/video RAM/priner/hard disk/removable media etc.) as well as software (processes/files/application program etc.)
 - Security: ensures isolation between processes and data of one user from other users. Also protects system programs as well as data on the system from external attackers.
-

OS: Views

- ❑ Extended machine
 - ❑ Virtual Machine
 - ❑ Resource manager
 - Multiplexes resources for efficiency
-

Evolution of OS

- ❑ Single user systems
 - ❑ OS – loader + library of common routines
 - ❑ Called supervisor or control program
 - Loads program and library functions from tape to memory
 - Transfers control to program
 - After program execution, copies CPU register values to core memory (core dumped). In case of an error, a programmer shall look at the dump to analyze what went wrong
-

Batch operating systems

- ❑ All programs written in one language to be run as a single batch
 - ❑ Shall reduce time for loading library tapes – done manually
 - ❑ OS = sequencer + loader + output program
-

Multiprogramming

- ❑ Running multiple programs concurrently
 - Keeping many programs in memory,
 - running one of them (say program A) and
 - switch to another (say program B) when I/O required for A
 - ❑ Aim was to improve resource utilization
 - Idle time of CPU should be minimized
 - ❑ Programs could belong to different users - multi-user
 - ❑ A multiuser OS is multiprogramming but not vice-versa.
-

Process Address Space

- ❑ How much memory is needed for an executable program?
 - Memory to be allocated for machine code
 - Memory to be allocated for all variables, constants and labels referred to in code
 - Memory required for code linked from libraries
 - Memory needed for maintaining
 - Stack (to facilitate function calls)
 - Heap (memory allocation at run time)
 - ❑ Requires program to be translated from high level language to machine language
-

Program Translation

❑ Compiler

- Translates HLL program into machine equivalent
- Creates table of incomplete instructions for unresolved addresses
 - **Library functions, Extern variables**

❑ Linker

- Static linking: Concatenates library with program and resolves call to library functions
- Dynamic linking: Address is resolved at runtime; OS maintains a table of DLL files loaded in the memory

❑ Loader

Program

- ❑ **Program:** a collection of machine instructions; each instruction is of format opcode [operand1, [operand2]]
 - ❑ **Identifier:** needs to be mapped to an address
 - Variable, constant, function name, label
 - ❑ **Statements:**
 - Data declaration: type of data (used by compiler to assess bytes required to store data)
 - Assignment
 - Expressions
 - Control constructs
 - Function calls
-

Statements

□ $X = Y$

- MOV accumulator, Y
- MOV X, accumulator

□ Expression: $A + B * C$

- MOV accumulator, B
 - MULT accumulator, C
 - ADD accumulator, A
-

Compiler Overheads

- ❑ Some variables are not defined by programmer but created by compiler during translation process.
 - Instruction labels
 - Temporary variables to store intermediate results (needed in expression parsing)
 - ❑ Compilers may also generate internal tables
 - Symbol table
 - VFT (virtual function table)
 - ❑ Memory required for storing data may be more than the sum of bytes required for data defined by the developer.
-

IF Statement

```
If (condition) {  
    Statement1  
    Statement2  
}  
Statement3
```

L1 and L2 are generated by compiler to ensure correct control flow.

Condition

Check condition == True

JNZ L2

Statement1

Statement2

L2: Statement3

IF..ELSE Statement

```
If (condition) {  
    Statement1  
}  
Else {  
    Statement2  
}  
Statement3
```

Condition

Check condition == True

JNZ L1

Statement1

JMP L2

L1: Statement2

L2: Statement3

While Statement

```
while (condition) {  
    Statement1  
}  
Statement2
```

L1:Condition

Check condition == True

JNZ L2

Statement1

JMP L1:

L2:Statement2

FOR Statement

```
for (init; cond; incr) {  
    Statement1  
}  
Statement2
```

```
init  
L1: cond  
Check cond == True  
JNZ L2  
Statement1  
incr  
JMP L1  
L2:Statement2
```


Assignment

```
If( cond1 )  
    Statement1  
Elseif(cond2)  
    Statement 2  
Else  
    Statement3  
Endif  
Statement4
```

```
Switch(value) {  
    case val1: stmt1; break;  
    case val2: stmt2; break;  
    case val3: stmt3; break;  
    default  
}
```

Switch Statement

```
Switch(expr) {  
  Case val1: stmt1; [break;]  
  Case val2: stmt2; [break;]  
  Case val3: stmt3; [break;]  
  ...  
  
  Default: stmtK;  
}  
Next;
```

```
V = expr  
If(V == val1) goto L1;  
Else if( V == val2) goto L2;  
...  
Else goto Lk;  
L1: stmt1; [goto next; ]  
L2: stmt2; [goto next; ]  
L3: stmt3; [goto next; ]  
  
Lk: stmtK;  
Next:
```


Assignment

Repeat
 Statement1
Until(cond)
Statement2

How would the machine
code translation change
when
break
continue
are used?


Source to Executable: Steps

- **Compilation:** translates a HLL program to Machine/Assembly language equivalent
- **Linking:** Resolves references to library functions
- **Loader:** Loads the program into memory
 - Absolute addresses are relocated
 - Relative addresses are retained




Source to Executable: Steps

- Compiled output is not machine executable
- Has unresolved references to undefined symbols
 - Library functions
 - Variables declared as **ext**
- **Linker resolves these addresses and stitches object files and library functions to create an executable image.**




Library



- ❑ A collection of files/functions
- ❑ Preamble is
 - a signature starting with magic bytes !<arch>
 - followed by a number of files
- ❑ Each file has a header followed by its contents
 - Header is fixed length
 - Data can be arbitrarily long


Static Linking



- ❑ Libraries added to each program
- ❑ Requires more memory space
- ❑ Multiple instances of same library in memory
- ❑ Faster as address need not be resolved at run-time


Dynamic Linking

- One instance of a library present in memory
- Library module loaded as needed
- All programs share libraries
- Call to library functions resolved at run time (dynamically linked library - DLL).
- OS keeps an import address table to locate each library module
 - Information on where a memory is loaded in memory
 - Used in resolving addresses at run-time



Source to Executable: Steps

- Linking: Resolves references to library functions
- Each library is collection of certain routines
- Multiple libraries may need to be linked
- Output contains executable file with extra information in headers
- Header information useful in loading program to memory
- Loader: Loads the program into memory
 - Absolute addresses are relocated
 - Relative addresses are retained



Address Space

- ❑ Compiler generates address space
 - ❑ Virtual address space – starts with zero
 - Virtual as it is not the physical address (where program is actually loaded)
 - Processes may have similar virtual address but are mapped to different physical address
 - ❑ Every instruction needs memory.
 - Number of bytes allocated for each instruction can vary
 - Even for same instruction, it could change depending on where operands are stored (Register/Memory): MOV
 - ❑ Every identifier needs an address
 - Variables
 - Constants
 - Function Names: Address of first instruction (CALL)
 - Labels: Address of an Instruction (GOTO) - assembler
-

Why is stack needed?

- ❑ Function A consists of n instructions.
 - ❑ At mth instruction, A calls another function B
 - ❑ Context should change from A to B
 - local variables of B should be referenced
 - IP should point to code in B
 - ❑ After the function B has executed
 - A should resume from where it left off?
 - ❑ If A calls itself recursively, new context of A should be created
 - ❑ Function calls: LIFO
-

Stack Frames

<pre>C() { }</pre>	<pre>B() { C(); A(0); }</pre>
<pre>A(n) { if(n is non-zero) B(); endif }</pre>	<pre>main() { A(1); }</pre>

main() Stack frame

Stack Frames

<pre>C() { }</pre>	<pre>B() { C(); A(0); }</pre>
<pre>A(n) { if(n is non-zero) B(); endif }</pre>	<pre>main() { A(1); }</pre>

main()
Stack frame
A
Stack frame

Stack Frames

<pre>C() { }</pre>	<pre>B() { C(); A(0); }</pre>
<pre>A(n) { if(n is non-zero) B(); endif }</pre>	<pre>main() { A(1); }</pre>

main() Stack frame
A Stack frame
B Stack frame

Stack Frames

<pre>C() { }</pre>	<pre>B() { C(); A(0); }</pre>
<pre>A(n) { if(n is non-zero) B(); endif }</pre>	<pre>main() { A(1); }</pre>

main() Stack frame
A Stack frame
B Stack frame
C Stack frame

Stack Frames

<pre>C() { }</pre>	<pre>B() { C(); A(0); }</pre>
<pre>A(n) { if(n is non-zero) B(); endif }</pre>	<pre>main() { A(1); }</pre>

main() Stack frame
A Stack frame
B Stack frame

Stack Frames

<pre>C() { }</pre>	<pre>B() { C(); A(0); }</pre>
<pre>A(n) { if(n is non-zero) B(); endif }</pre>	<pre>main() { A(1); }</pre>

main() Stack frame
A Stack frame
B Stack frame
A Stack frame

Stack Frames

<pre>C() { }</pre>	<pre>B() { C(); A(0); }</pre>
<pre>A(n) { if(n is non-zero) B(); endif }</pre>	<pre>main() { A(1); }</pre>

main() Stack frame
A Stack frame
B Stack frame

Stack Frames

<pre>C() { }</pre>	<pre>B() { C(); A(0); }</pre>
<pre>A(n) { if(n is non-zero) B(); endif }</pre>	<pre>main() { A(1); }</pre>

main()
Stack frame
A
Stack frame

Stack Frames

<pre>C() { }</pre>	<pre>B() { C(); A(0); }</pre>
<pre>A(n) { if(n is non-zero) B(); endif }</pre>	<pre>main() { A(1); }</pre>

main() Stack frame

Stack Layout: Activation Frame

- ❑ What do we do when we call a function?
 - What data need to be stored?
 - Where do they go?

 - ❑ How do we return from a function?
 - What data need to be restored?
 - Where do they come from?
-

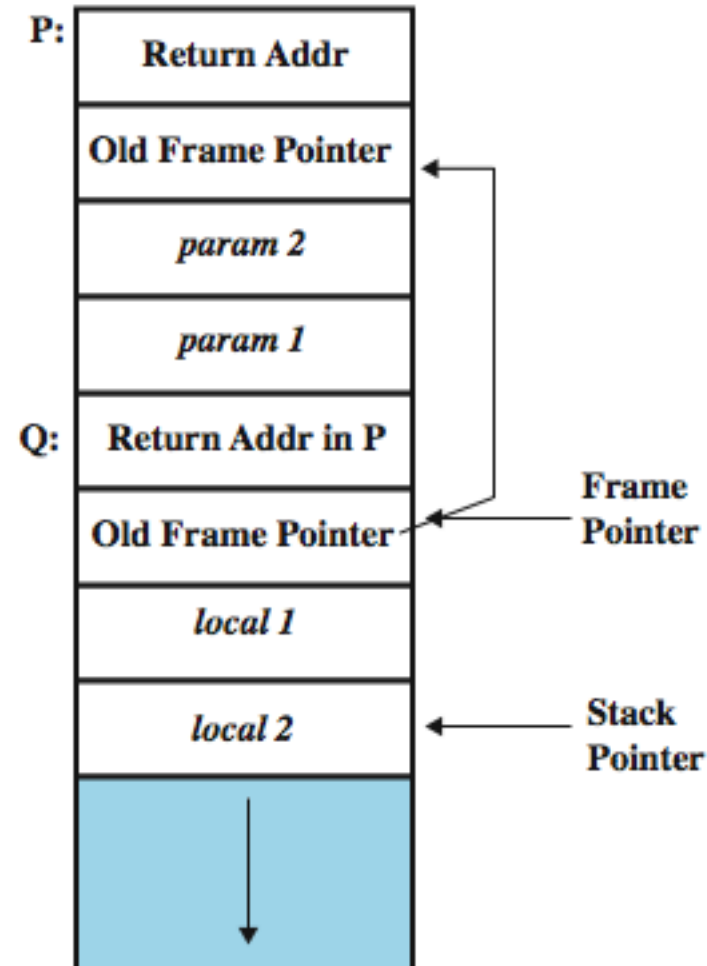
Activation Frame

- ❑ On each function call, an activation record or activation record is pushed onto stack
 - ❑ Activation stack/frame consists of
 - parameters to be passed to the called function,
 - return address in callee function (instruction address where control shall return to after the called function has been executed),
 - old stack pointer value local variables of the called function,
-

Activation Frames (contd).

Calling function: needs a data structure to store the “return” address and parameters to be passed

Called function: needs a place to store its local variables (different for every call)



Creating New Activation Frame

❑ Calling function:

- Push arguments onto the stack (in reverse order)
- Push the address of the instruction to run after control returns to you
- Jump to the function

❑ Called function:

- Push the old frame pointer onto the stack (%ebp)
 - Set current frame pointer (%ebp) to where the end of the stack is right now (%esp)
 - Push local variables onto the stack
-

Activation Frame

Previous Stack Frame (main)Di
Previous Stack Frame (F1)
Function Arguments to F2 (in reverse) rightmost argument first)
Return Address (instruction address in F1)
<Previous Frame Pointer: Contents of ebp register>
Local Variables
Local Buffer Variables

- ❑ Arguments in reverse order
- ❑ Return address
- ❑ Store old frame pointer (%ebp) - so that context can be returned to callee function after this call
- ❑ Set current frame pointer (%ebp) to where the end of the stack is right now (%esp)
- ❑ Push local variables onto the stack

Different green shades show how calling and called functions contribute to activation record or stack frame. For more details, visit

<https://www.cs.princeton.edu/courses/archive/spr03/cs320/notes/7-1.pdf>

<https://softwareengineering.stackexchange.com/questions/195385/understanding-stack-frame-of-function-call-in-c-c>

Removal of Activation Frame

- ❑ Called function (to return):
 - Deallocate local variables: `%esp = %ebp`
 - Restore base pointer: `pop %ebp`
 - Jump back to where they wanted us to: `%eip = (%esp)`
 - ❑ Calling function (on return):
 - Remove arguments from stack
-

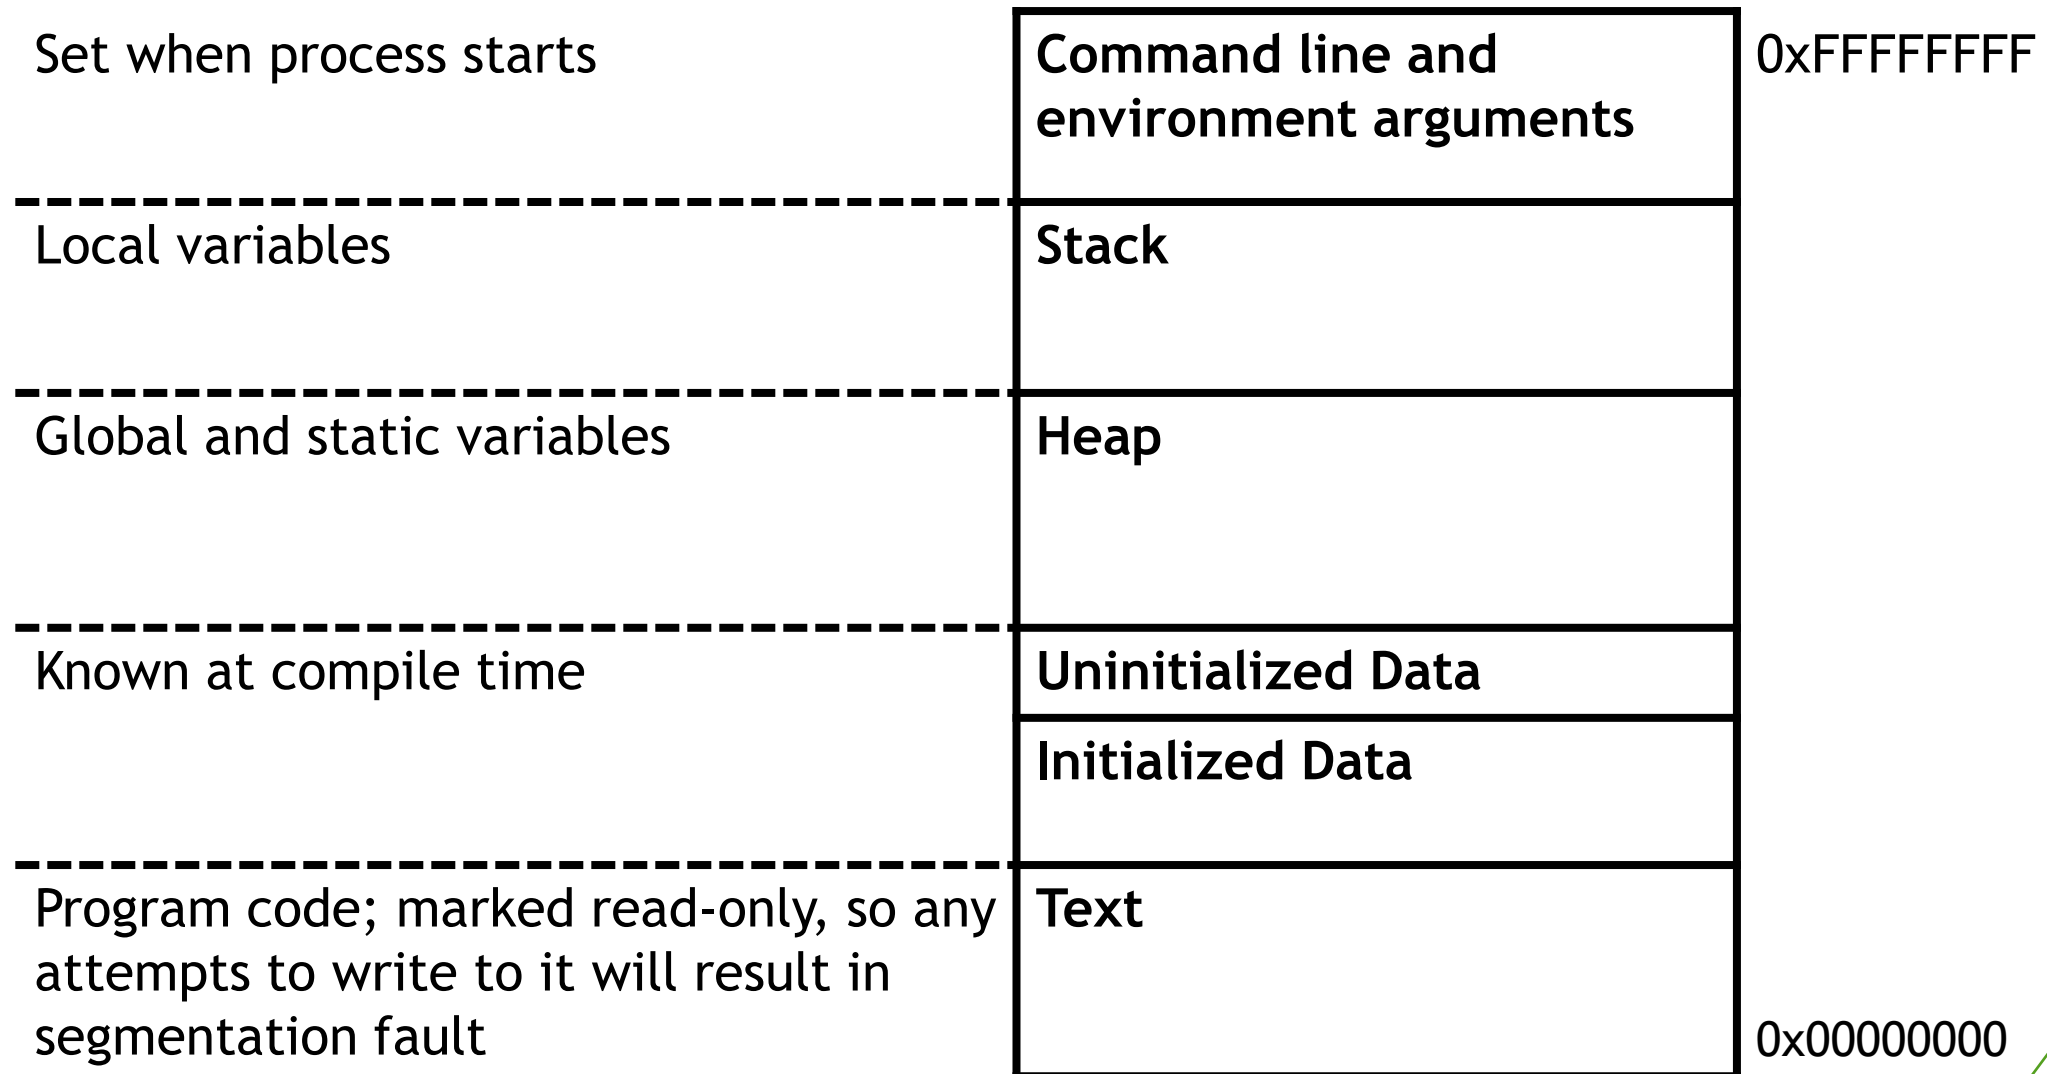
Loader

- ❑ Absolute: Copies program to memory starting from a specific address (say 5000H)
 - ❑ What if the program needs to be mapped to another memory location? (say 7000H)
 - All identifier addresses need to be incremented by 2000H
 - But what about relative addresses? JMP offset
 - ❑ Relocatable Loader: modifies the object program by translating addresses such that the program can be loaded at an address different from the location originally specified
-

Process

- ❑ Process – a program in execution; process execution must progress in sequential fashion
 - ❑ Process address space consists of
 - Text section: the program code
 - Current activity including program counter, processor registers
 - Stack containing temporary data, Function parameters, return addresses, local variables
 - Data section containing global variables
 - Heap for memory dynamically allocated
-

Process Virtual Address Space

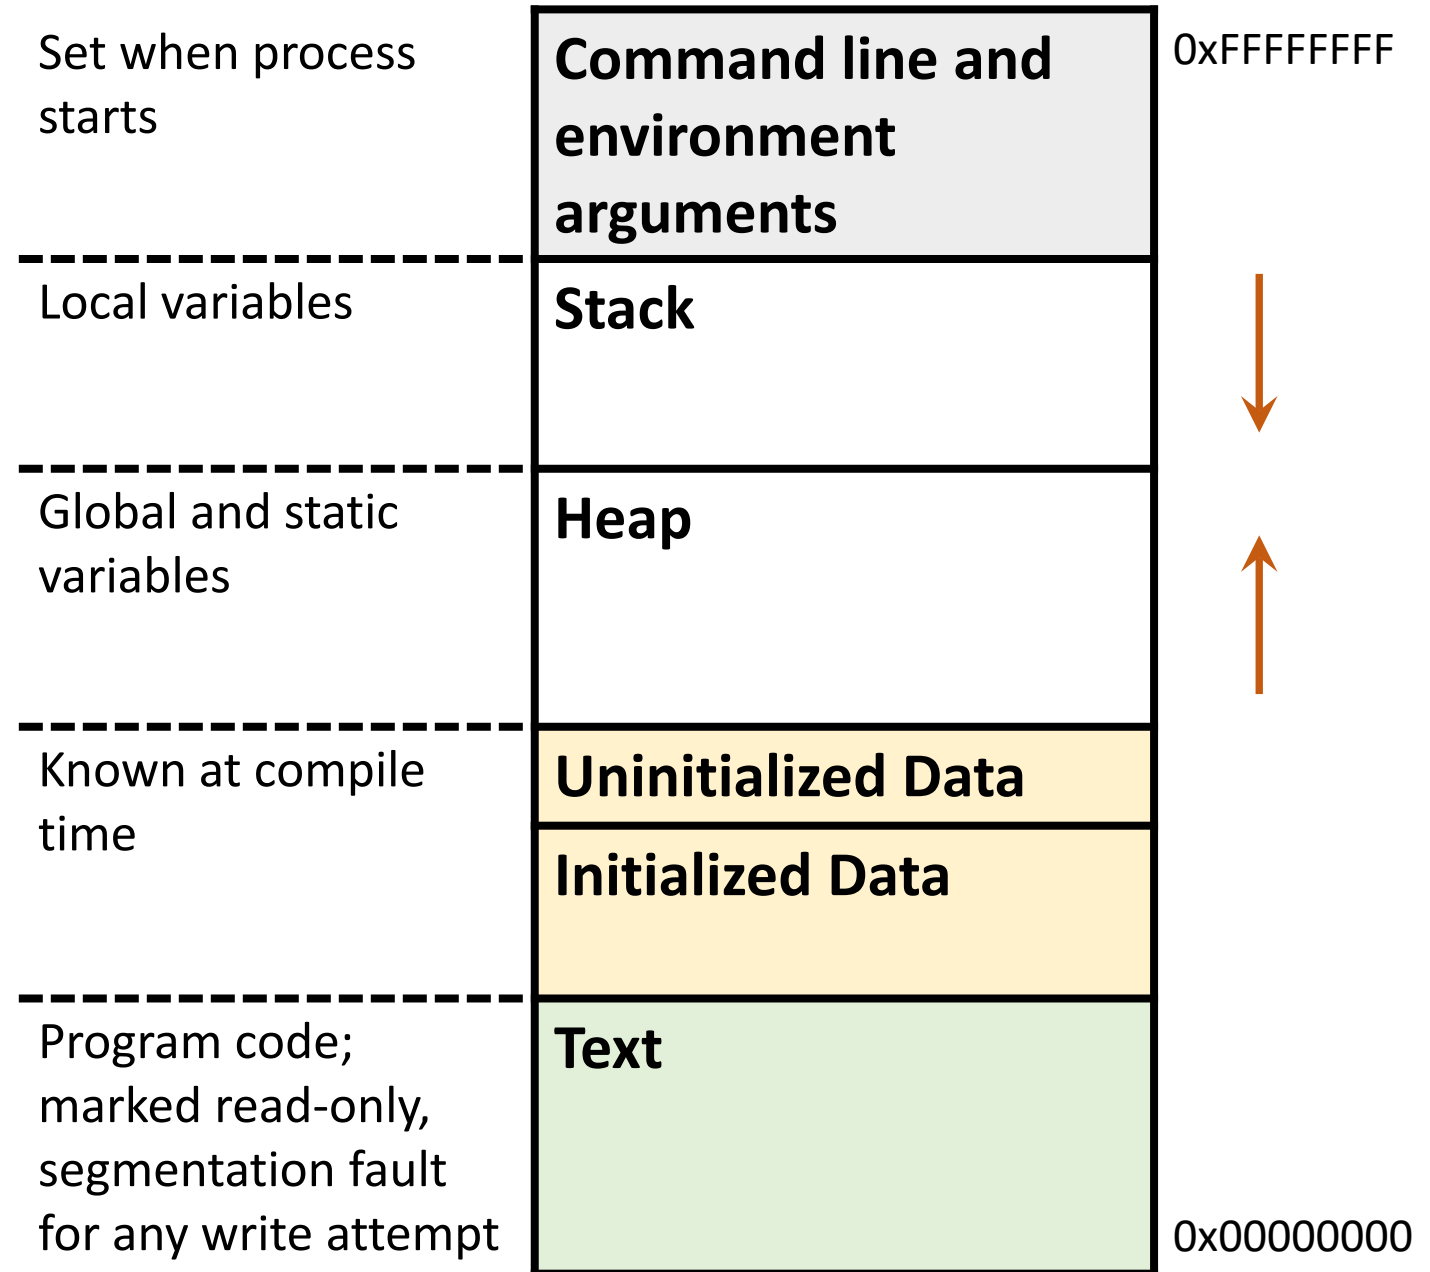


Address space

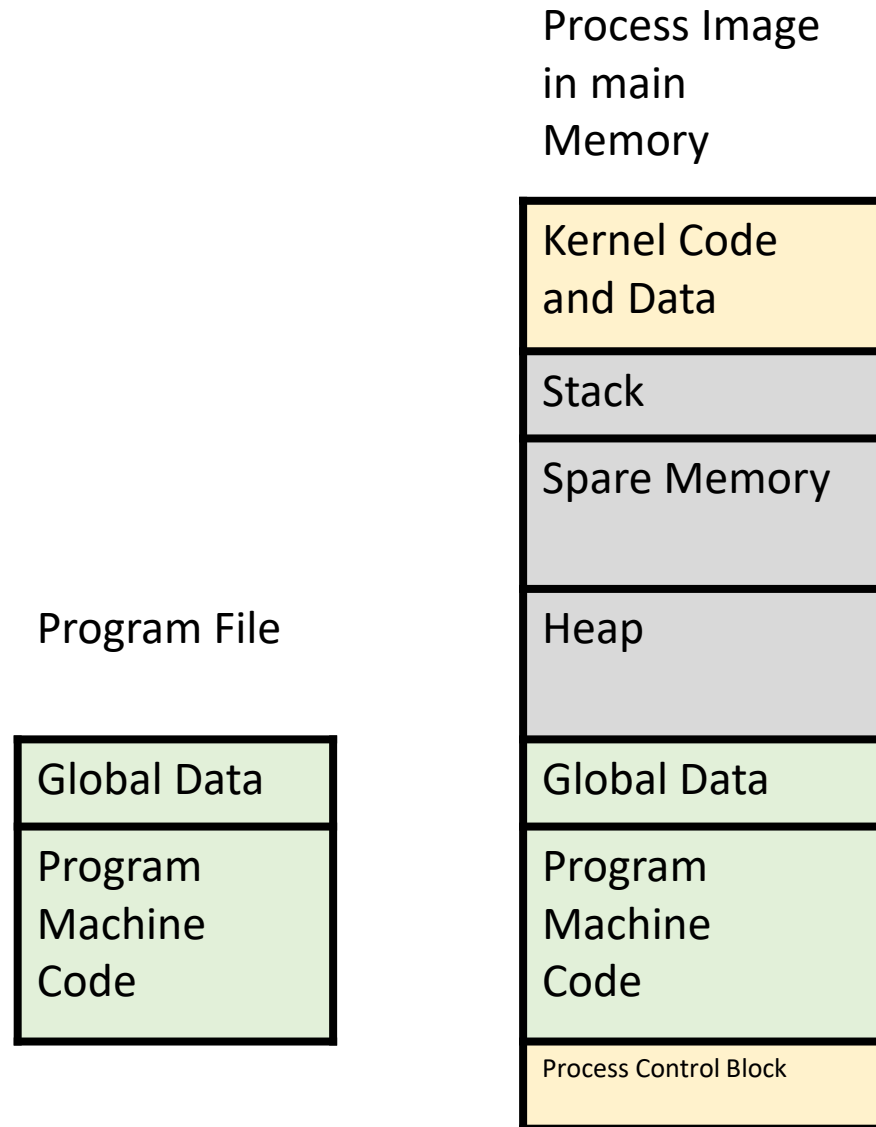
- for code instructions
 - Code Segment (Text)
- for Global and Static data – defined at compile time
 - Data Segment (BSS)
- for dynamic data – allocated at run-time through `malloc()`, `new()`
 - Heap
- Virtual address space always starts from Zero
 - Stack Segment

https://en.wikipedia.org/wiki/X86_memory_segmentation

Stack and heap grow in opposite directions



Programs and Processes



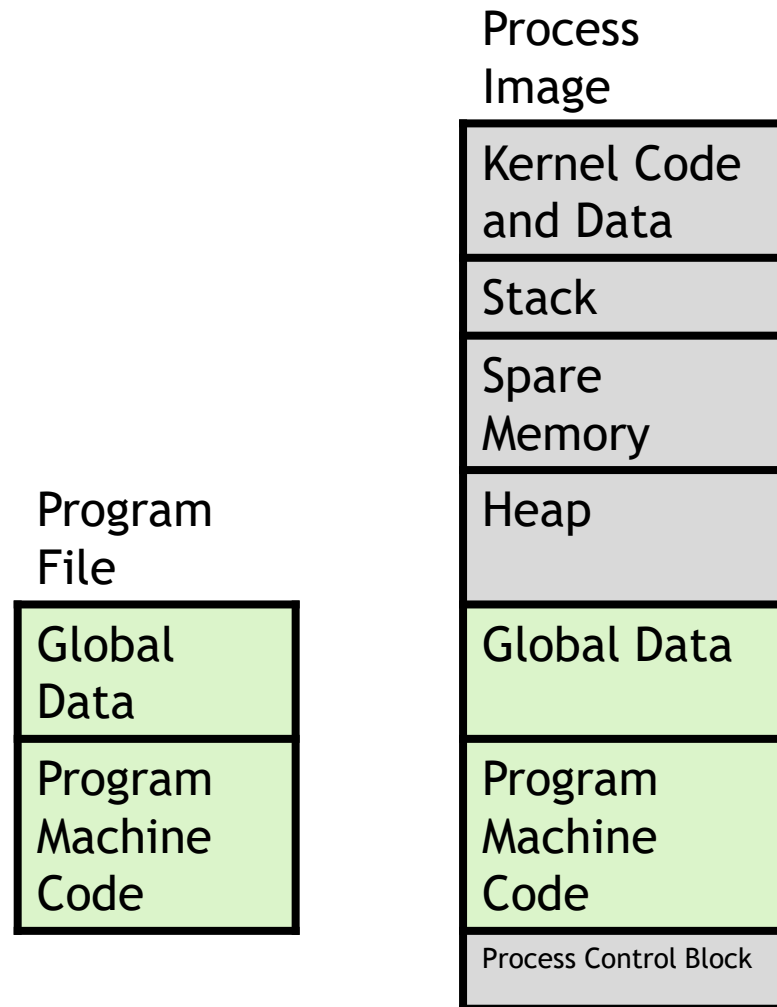
Process Address Space

- ❑ Starting address of executable program is 0000H.
 - ❑ Virtual address space
 - ❑ Physical address space: the memory location where program is actually loaded
 - ❑ Memory Management Unit: Translates from virtual address to physical address
 - ❑ Memory management
 - Paging
 - Segmentation
-

Context of a Process

- ❑ In multiprogramming environment, CPU time is shared between different processes.
 - ❑ At time $t = t_0$, process A running.
 - ❑ After a quantum (τ , $t = t_0 + \tau$), another process B is scheduled to run on CPU.
 - ❑ Status of A needs to be saved/preserved so that it can be resumed from same point
 - Contents of PC and other CPU registers
 - File descriptors opened by A; resources allocated
 - Accounting information
 - ❑ OS needs to maintain data structures to keep this information (context).
-

Program v/s Process



- In multi-processing environment, CPU time is shared between different processes.
- When CPU time allotted to a process, say A, is exhausted or process needs I/O, CPU is allocated to another process, say B.
- OS needs to ensure that when A is given back CPU, it resumes from where it left.
- For this, OS needs information (process)
 - Context: Contents of all registers including PC
 - Resources allocated to process
- This extra housekeeping information (dynamic) is what makes a process different from program.

Process Control Block

- ❑ PCB: stores all the information about a process.
 - ❑ PCB is also called process descriptor.
 - ❑ PCB is updated during as a process changes its state.
 - ❑ One PCB per process.
-

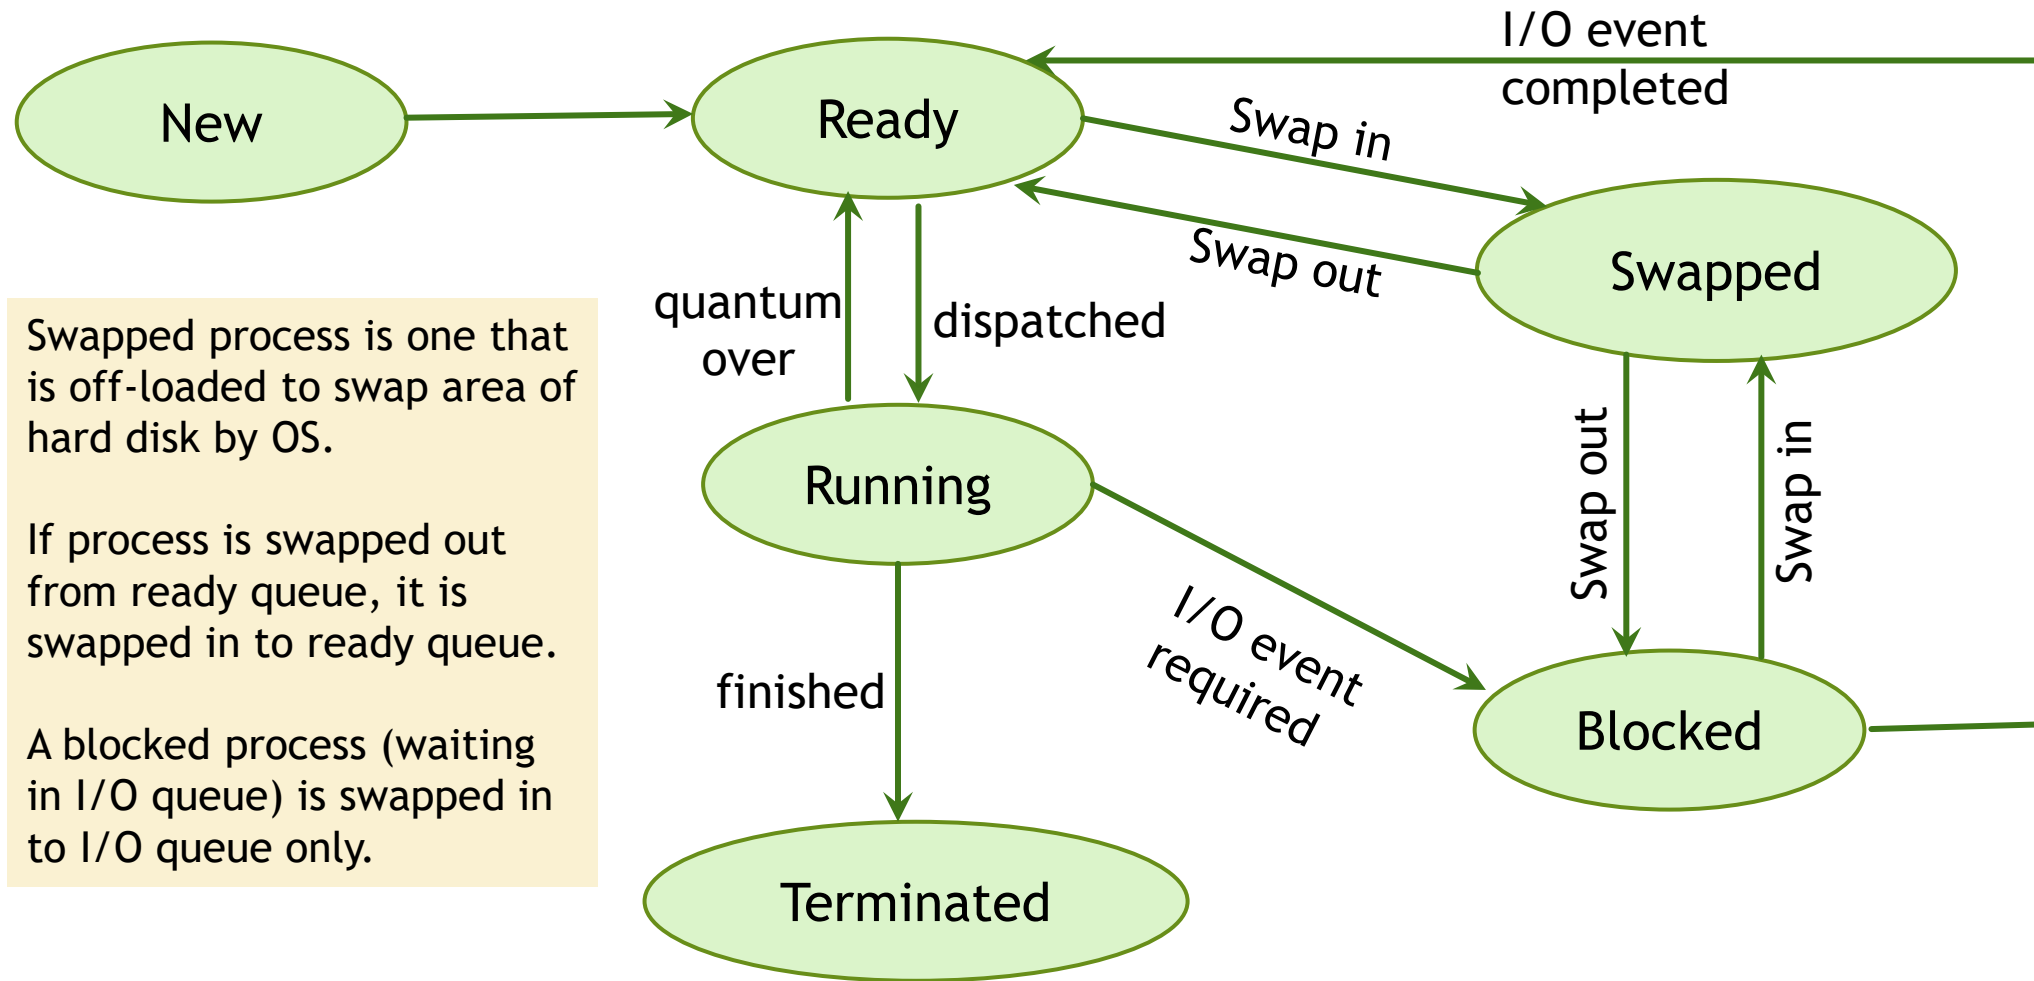
Process Control Block

- ❑ Process ID, owner, Privileges
 - ❑ Process state – “ready”, “running”, "suspended",
 - ❑ Child process IDs, interacting processes
 - ❑ IPC information –flags, signals and messages
 - ❑ Program Counter (PC) and other CPU registers
 - ❑ Scheduling Information – priority, pointers to scheduling queue
 - ❑ Memory Management Information–page table, memory limits, segment table
 - ❑ Accounting Information–amount of CPU used for process execution, time limits, execution ID etc.
 - ❑ I/O Status Information: I/O devices allocated, file descriptors
-

Process States

- ❑ New: request to run a program received; loader loads program image to memory; OS creates a new PCB; enters process ID in process table
 - ❑ Ready: process waiting in queue to be scheduled
 - ❑ Running: process allocated CPU
 - ❑ Blocked/Waiting: process waiting for I/O event (put in respective queue)
 - ❑ Swapped: process copied to hard disk (swap area)
 - ❑ Terminated: Process has finished execution
-

Process State Diagram



Swapped process is one that is off-loaded to swap area of hard disk by OS.

If process is swapped out from ready queue, it is swapped in to ready queue.

A blocked process (waiting in I/O queue) is swapped in to I/O queue only.

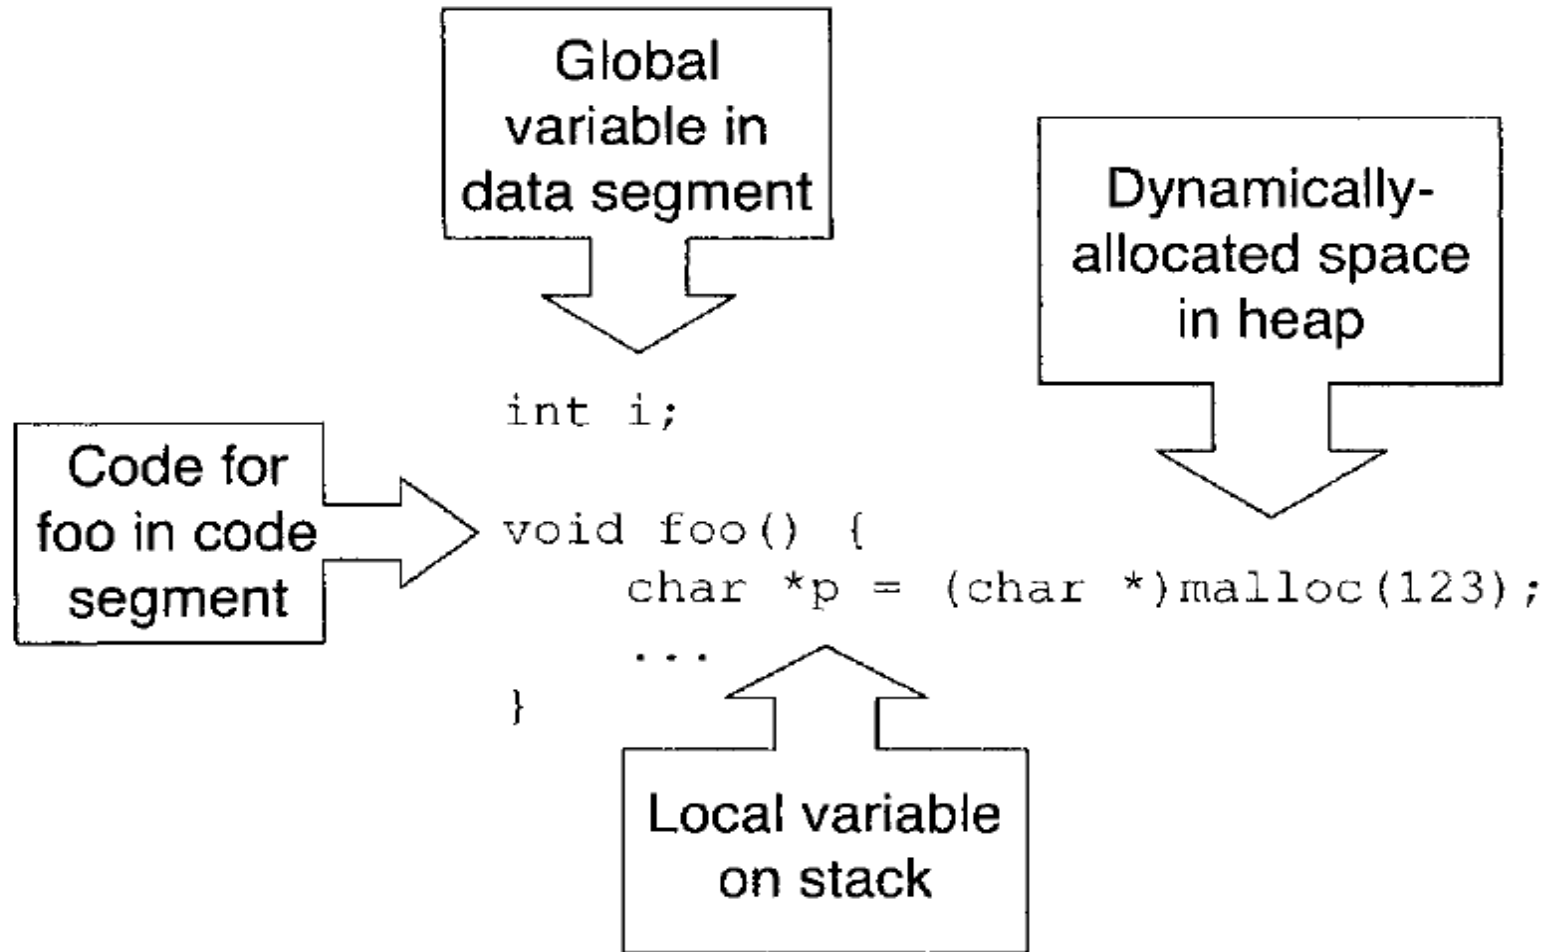
Processes and Threads

- ❑ To increase throughput of a process, threads are used.
 - for a certain functionality of the process.
 - When one thread is blocked, others can continue; the process need not be blocked.
 - Process can complete its quanta; context switches among process can reduce.
 - ❑ Threads are independent sequences of instructions that are executed by the CPU without waiting for other threads.
 - ❑ Threads within a process all share the same memory space, but each has its own context (processor registers) and stack.
 - ❑ Any switches between threads results in saving all values in the CPU are saved in a structure called the thread context. The OS then loads the thread context of a new thread into the CPU and executes the new thread.
-

Segments in x86 Architecture

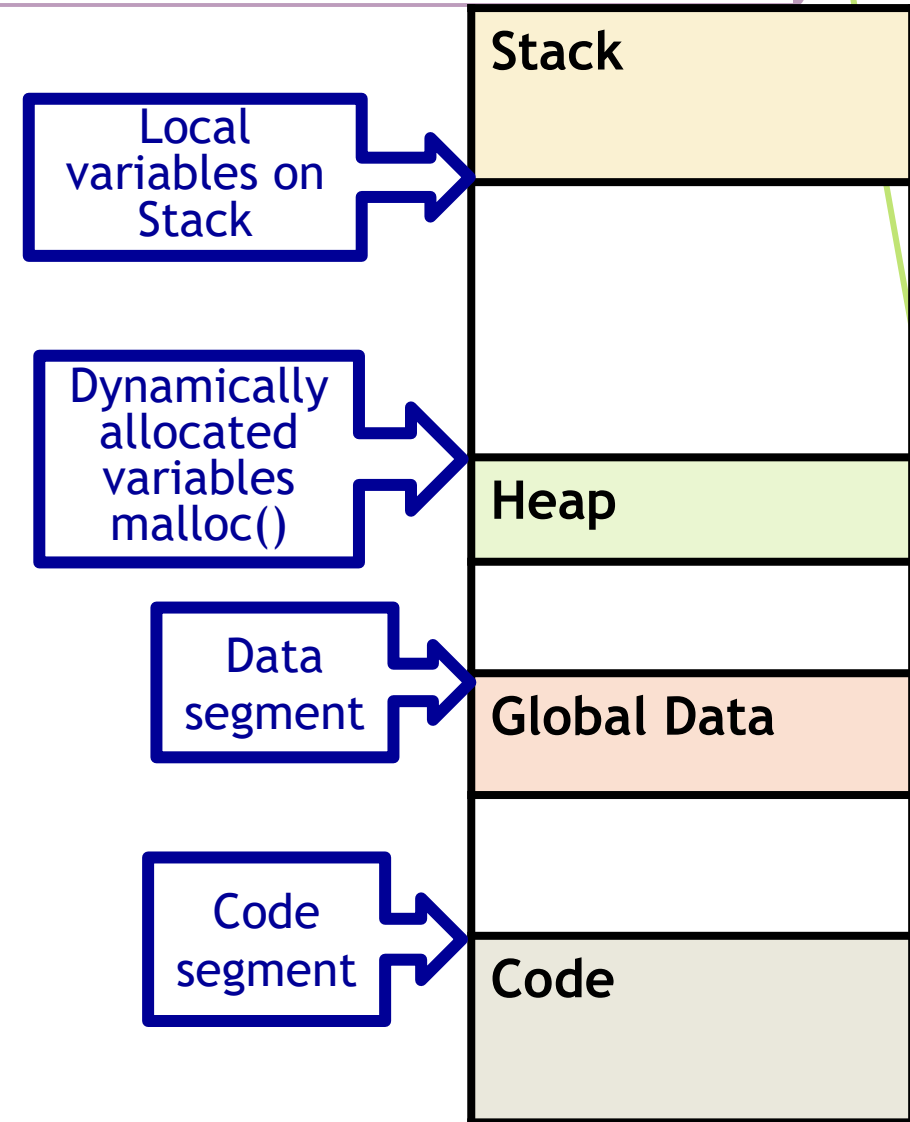
- ❑ Multi-programming requires isolation between processes
 - ❑ Limit registers used for enacting isolation
-

Segments in x86 Architecture



Segments in x86 Architecture

- ❑ Multi-programming requires isolation between processes
 - Limit registers
 - Upper and lower bound for address
- ❑ One way is to use segments
 - $\text{Effective Address} = [\text{Segment Register}] + \text{Address}$
 - Contiguous memory not needed
- ❑ Code segment: instructions of the program code
- ❑ Stack segment: stack
- ❑ Data segment: global data
- ❑ Heap: Dynamically allocated data



Processes v/s Threads

- ❑ A process contains one or more threads that are executed by the CPU.
 - ❑ Processes also contribute to stability by preventing errors or crashes in one program from affecting other programs.
 - ❑ Processes are the container for execution, but threads are what the OS executes.
 - ❑ Threads are independent sequences of instructions executed by the CPU without waiting for other threads.
 - ❑ A process contains one or more threads, which execute part of the code within a process.
 - ❑ Threads within a process all share the same memory space, but each has its own processor registers and stack.
 - ❑ Thread Context: When one thread is running, it has complete control of the CPU, or the CPU core, and other threads cannot affect the state of the CPU or core.
 - ❑ When a thread changes the value of a register in a CPU, it does not affect any other threads. Before an OS switches between threads, all values in the CPU are saved in a structure called the thread context. The OS then loads the thread context of a new thread into the CPU and executes the new thread.
-



Thank you.

