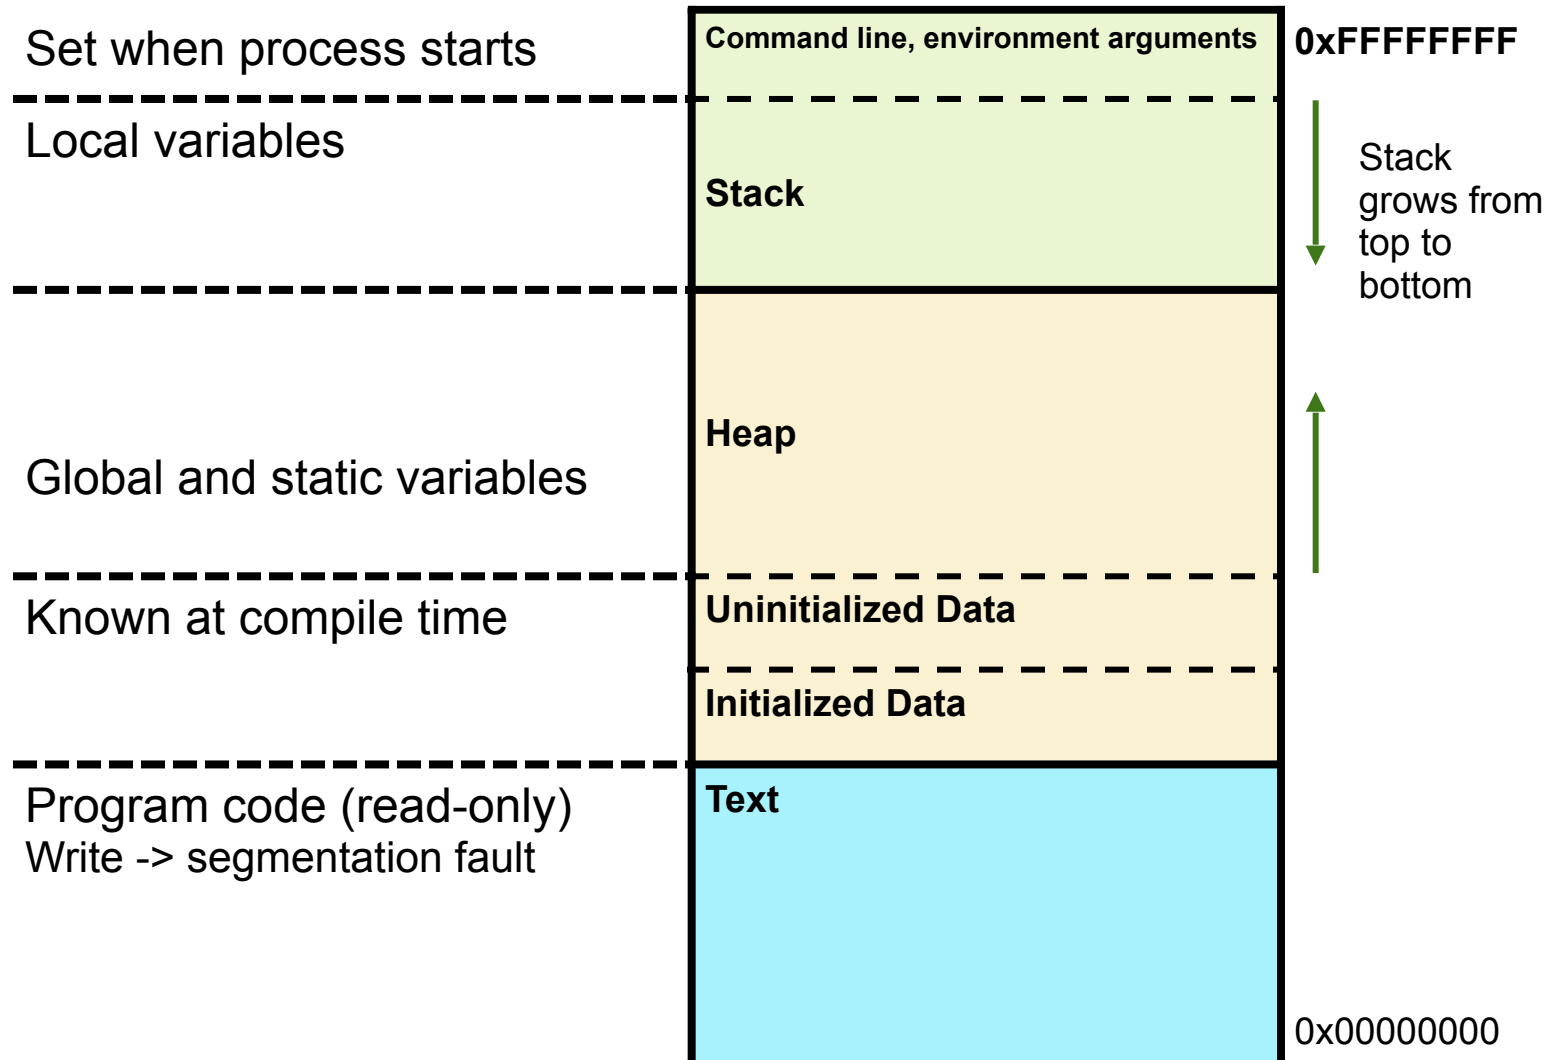


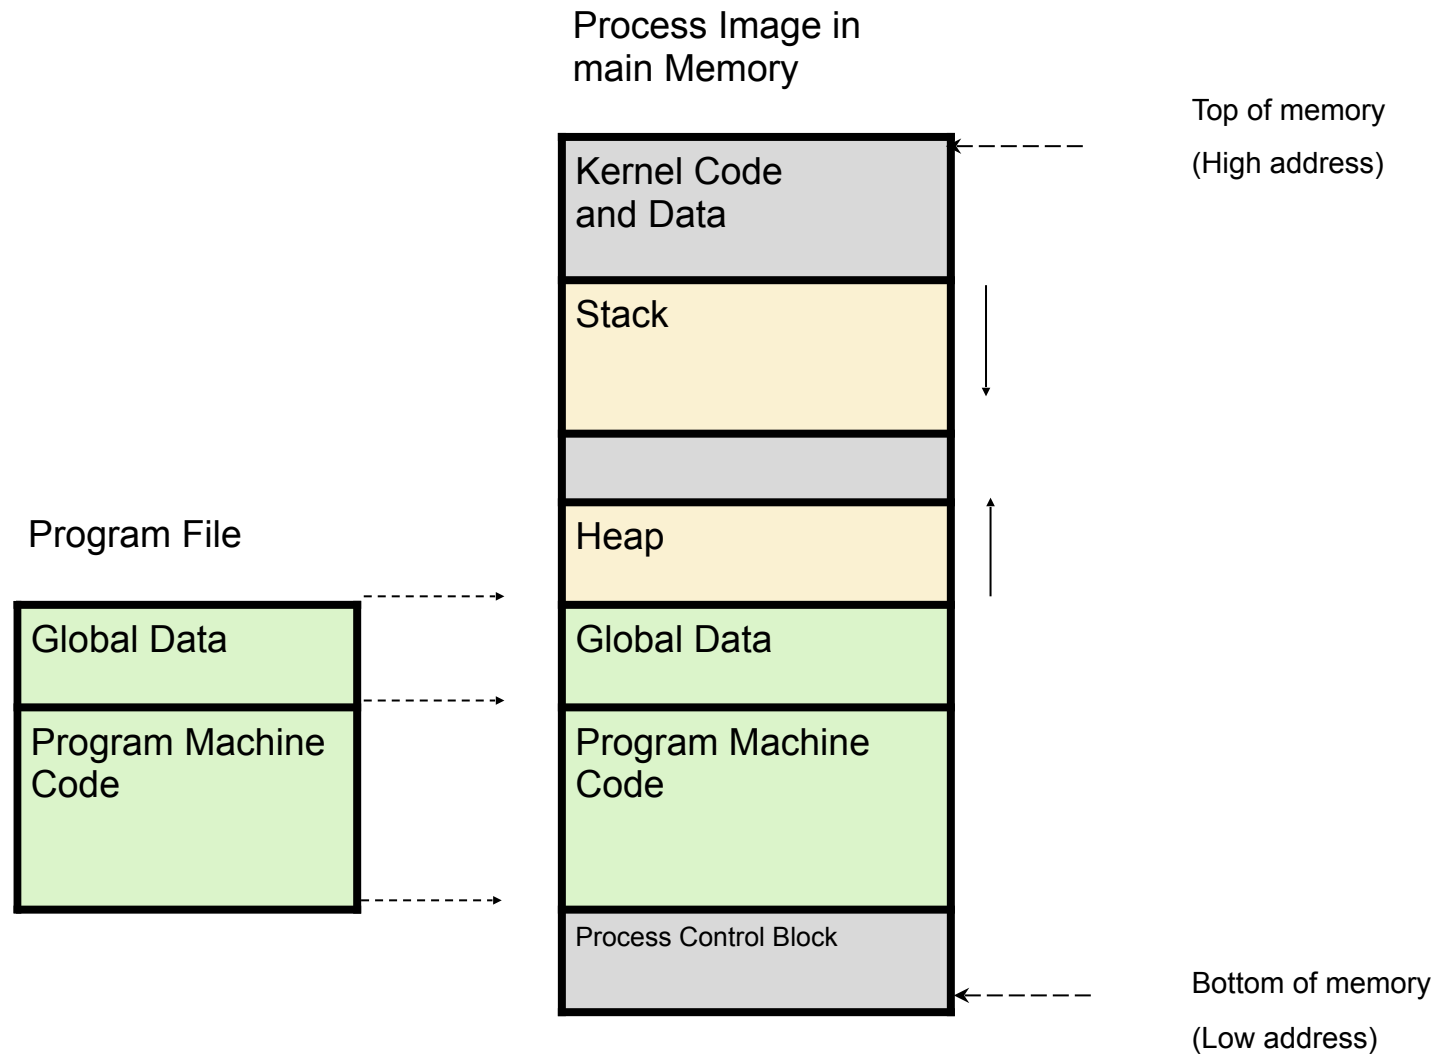
Memory Management - II

Segmentation and Paging

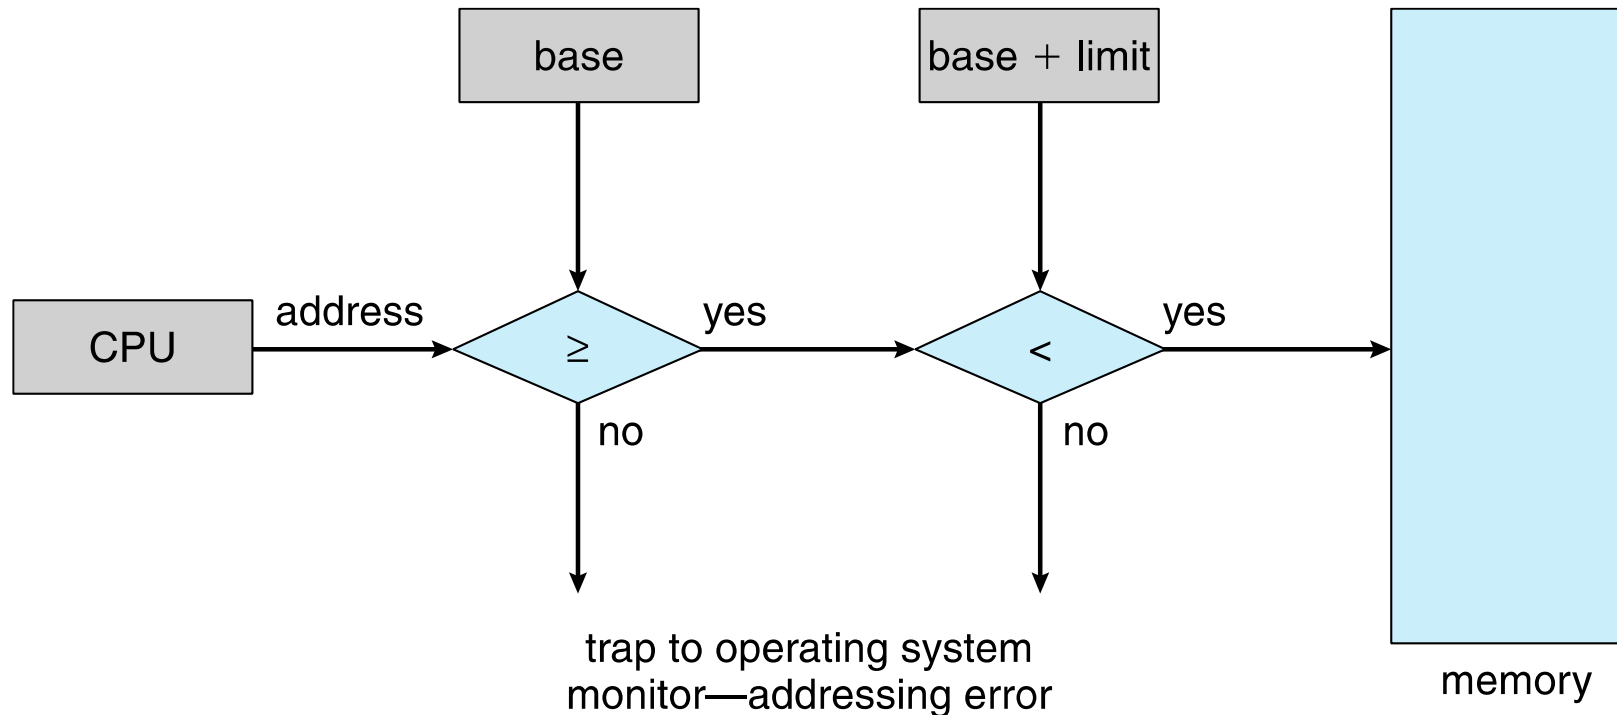
Process Virtual Address Space



Address Space: Processes



Memory Protection



Process Address Space

- ❑ Virtual address space:
 - Starting address is 0000H.
 - Includes space for heap and stack
 - ❑ Physical address space: RAM region actually allocated
 - ❑ Memory Management Unit
 - Translates virtual address to physical address
 - Checks that memory violation does not occur
 - ❑ Memory management
 - Paging
 - Segmentation
-

Memory Virtualization

- ❑ Space Sharing

- ❑ Hardware Support

- Bound registers – one pair (base and limit) per CPU core
 - Segments and CPU registers (x86)
 - Code Segment
 - Stack Segment
 - Data Segment
 - Extra Data Segment
 - Context switch: registers saved (outgoing) and restored (incoming)
-

Segmentation

- ❑ Three logical segments
 - Code, stack and heap
 - Two bits needed to represent each segment type
 - ❑ Virtual address: $\langle \textit{Segment type}, \textit{Offset} \rangle$
 - ❑ Segment register: 16 bits
 - 2 bits reserved for type [00: code, 01: stack, 10: heap]
 - 14 bits for offset
 - ❑ Segment size = $2^{14} = 16 \text{ KB}$
 - ❑ Code of program limited to 16 KB
 - ❑ Size of program limited to 48 KB
-

Segmentation

- ❑ One base, limit register pair per logical segment
 - code, stack, heap
 - ❑ Allocation:
 - Contiguous memory allocated to each segment
 - Segments can be allocated independently
 - ❑ Address represented as base and offset
 - Base address of segment stored in register
 - Offset: number of bits (say m)
 - Offset determined by Segment Size
-

Segmentation : Memory Allocation

Segment 1
Segment 2
Segment 3
Segment 4

Virtual Address Space

Physical Memory

Segment 3
Segment 2
Segment 1
Segment 4

Segmentation: Hardware Support

- ❑ Logical/Virtual address consists <segment-number, offset>,
 - ❑ Segment table – maps virtual address to physical address
 - base – starting physical address of segment
 - limit –length of the segment
 - ❑ Segment-table base register (STBR): location of segment table
 - ❑ Segment-table length register (STLR): number of segments
 - ❑ For valid address
 - Segment number < STLR
 - Offset < limit
-

Segmentation: Address Translation

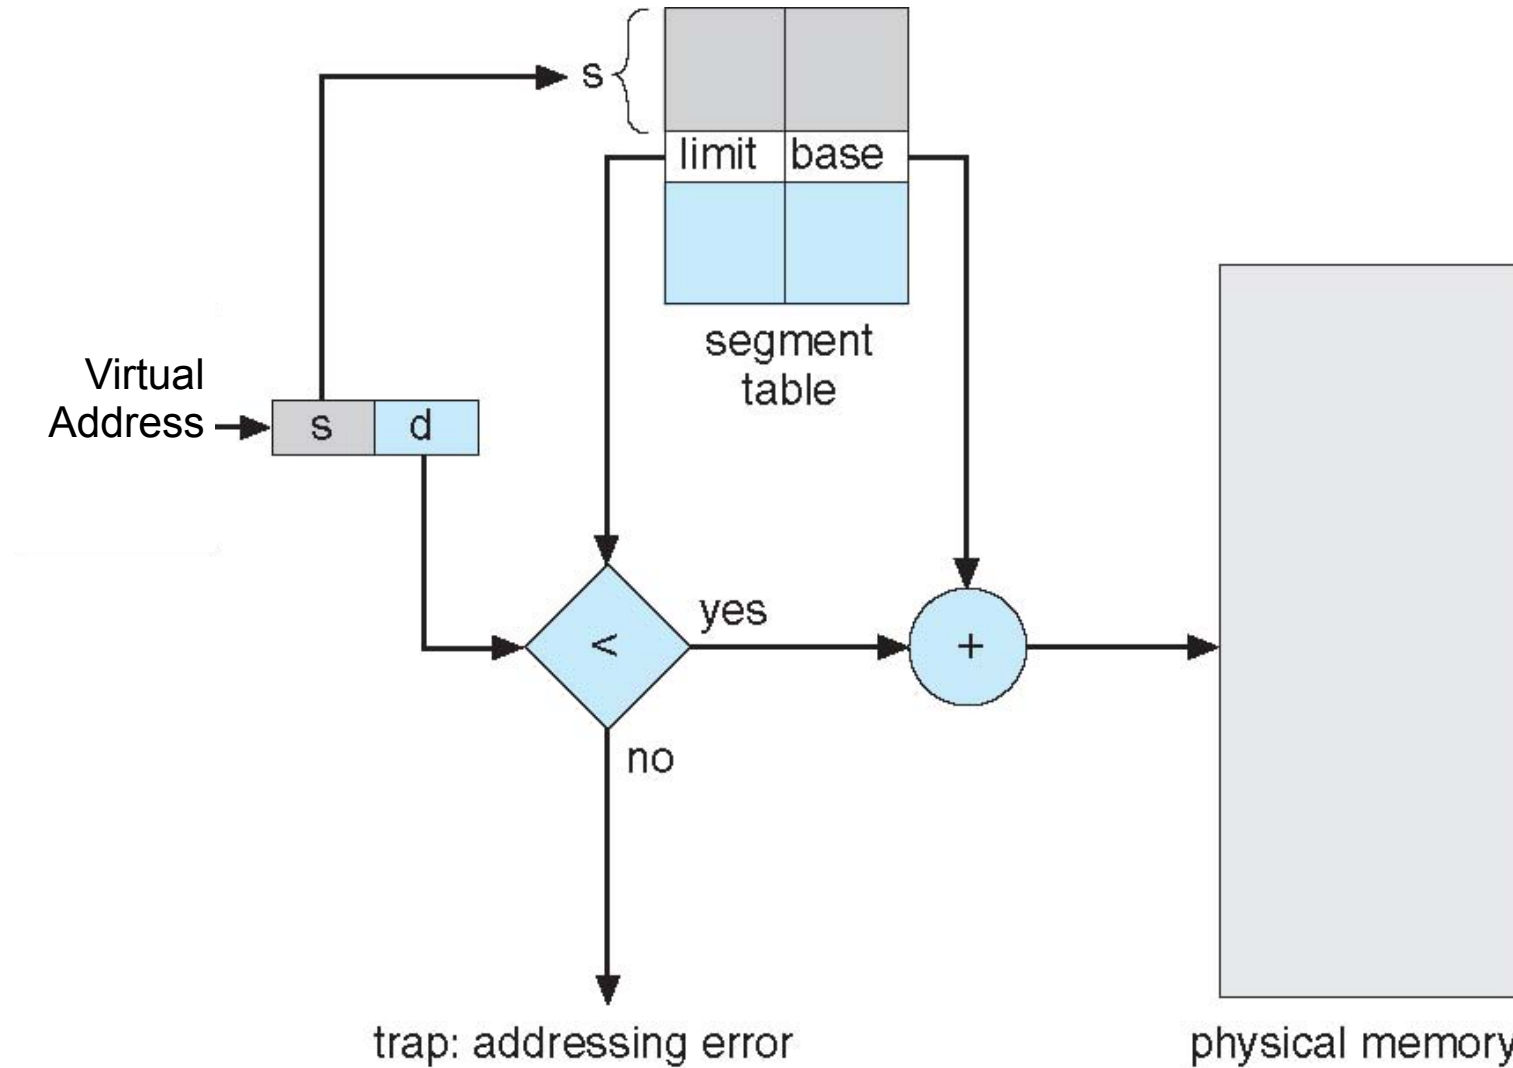
❑ For valid address

- Segment number < STLR
- Offset < limit


❑ Address translation:

- Use segment number as index into Segment Table
 - Base address = ST[segment number]
 - If (offset < limit)
 - **Physical address = Base Address + offset**
 - Else
 - **Segmentation Fault**
 - Endif
-

Segmentation: Address Translation




Segmentation Address Translation: Example



- ❑ Virtual address: 16 bits
- ❑ Segment size: 1 KB = 0400 H
- ❑ Base address: 4000 H
- ❑ Limit register: 43FF H
- ❑ Virtual address: 1356 H
- ❑ Offset: $(1356 - 1240) \text{ H} = 0116 \text{ H}$
- ❑ Physical address: $(4000 + 0116) = 4116 \text{ H}$
- ❑ $(4116 > 4000)$ Yes
- ❑ $(4116 < 43FF)$ Yes
- ❑ Access allowed: Yes

Segmentation Address Translation: Example

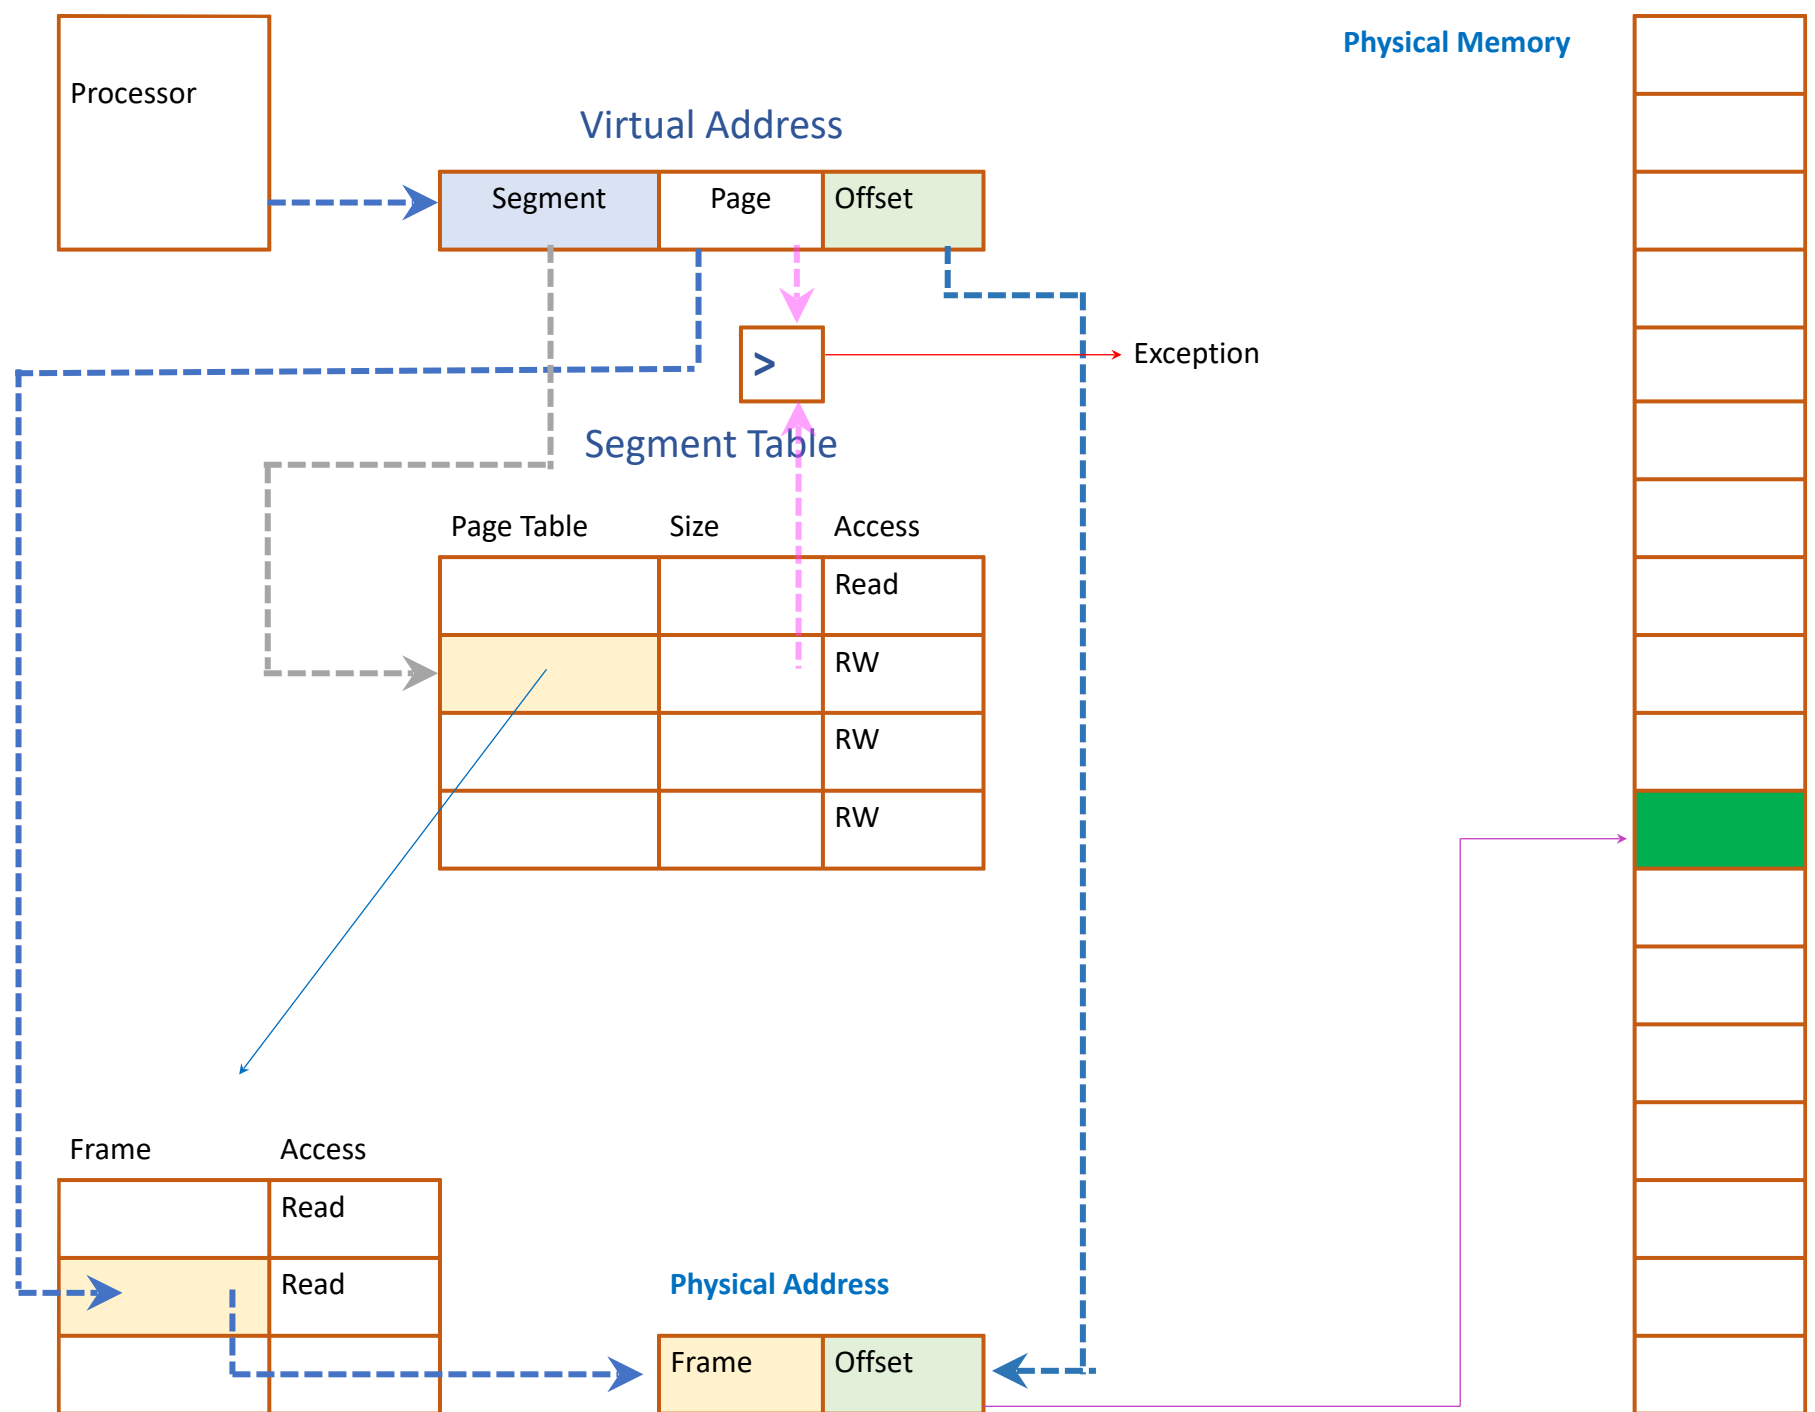


- ❑ Virtual address: 16 bits
- ❑ Segment size: 1 KB = 0400 H
- ❑ Base address: 4000 H
- ❑ Limit register: 43FF H
- ❑ Virtual address: 1736 H
- ❑ Offset: $(1736 - 1240) \text{ H} = 04F6 \text{ H}$
- ❑ Physical address: $(4000 + 0456) = 4456 \text{ H}$
- ❑ $(4456 > 4000) : \text{Yes}$
- ❑ $(4456 < 43FF) : \text{No}$
- ❑ **Access allowed : No**

Challenges with Segmentation

- ❑ Logical entities: Variable size
 - ❑ Segment size needs to be fixed; limited by size of base and limit registers
 - ❑ Otherwise, segmentation needs to be paged
 - Increases complexity
 - Increases access time
-

Paged Segmentation



Paging

- ❑ Virtual address space divided into equal sized pages (non-overlapping but contiguous)
 - ❑ Physical memory also divided into same sized blocks (called frames)
 - ❑ Page size = 2^m bytes
 - ❑ OS loads pages to physical frames
 - creates page table per process to store mapping
 - PTBR (page table base register): address of page table
 - Page table entries (PTE) changed as pages are loaded to/removed from main memory
 - At a given time, only a subset of process pages may be loaded
 - ❑ Internal fragmentation (average) = half a page size
-

Paging: Address Translation

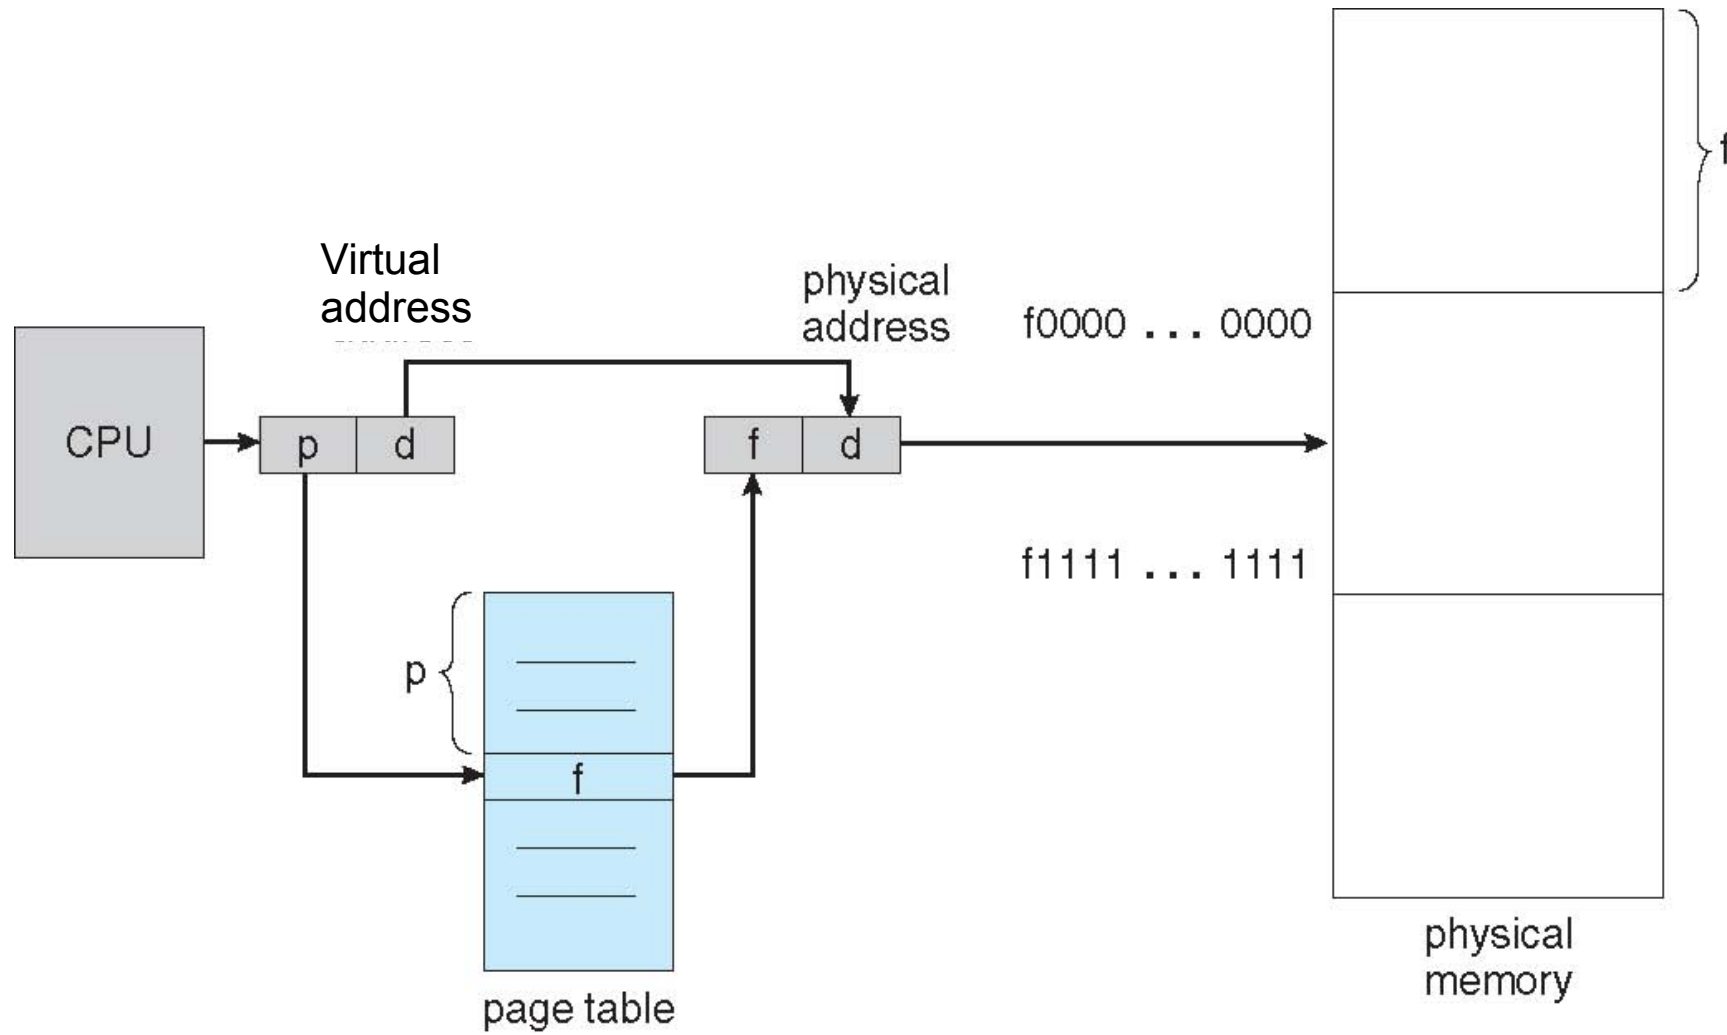
- ❑ Virtual address is a 2-tuple $\langle vp, d \rangle$
 - Virtual page number (vp) – index into page table that contains base address of each page in physical memory
 - Page offset (d) – added to base address of the page to obtain the physical memory address
- ❑ Virtual address of m -bits (Virtual space = 2^m bytes)
- ❑ Page size = 2^n bytes (n -bits)
- ❑ Physical address is a 2-tuple $\langle p, d \rangle$
- ❑ Address translation: replacing vp by $p = PT[vp]$

bits

bits

Virtual Page Number (VPN)	Page Offset
---------------------------	-------------

Paging: Address Translation



Linux: Process Memory Map

- ❑ How a process is mapped to memory

cat /proc/pid/maps

- ❑ List of all the memory mappings for the process

cat /proc/pid/pagemap

- ❑ Details: <https://kernel.org/doc/html/latest/admin-guide/mm/pagemap.html>
-

Windows: Process Memory Map


- ❑ List of all the memory mappings for the process with the specified PID.

tasklist /v /pid <pid>

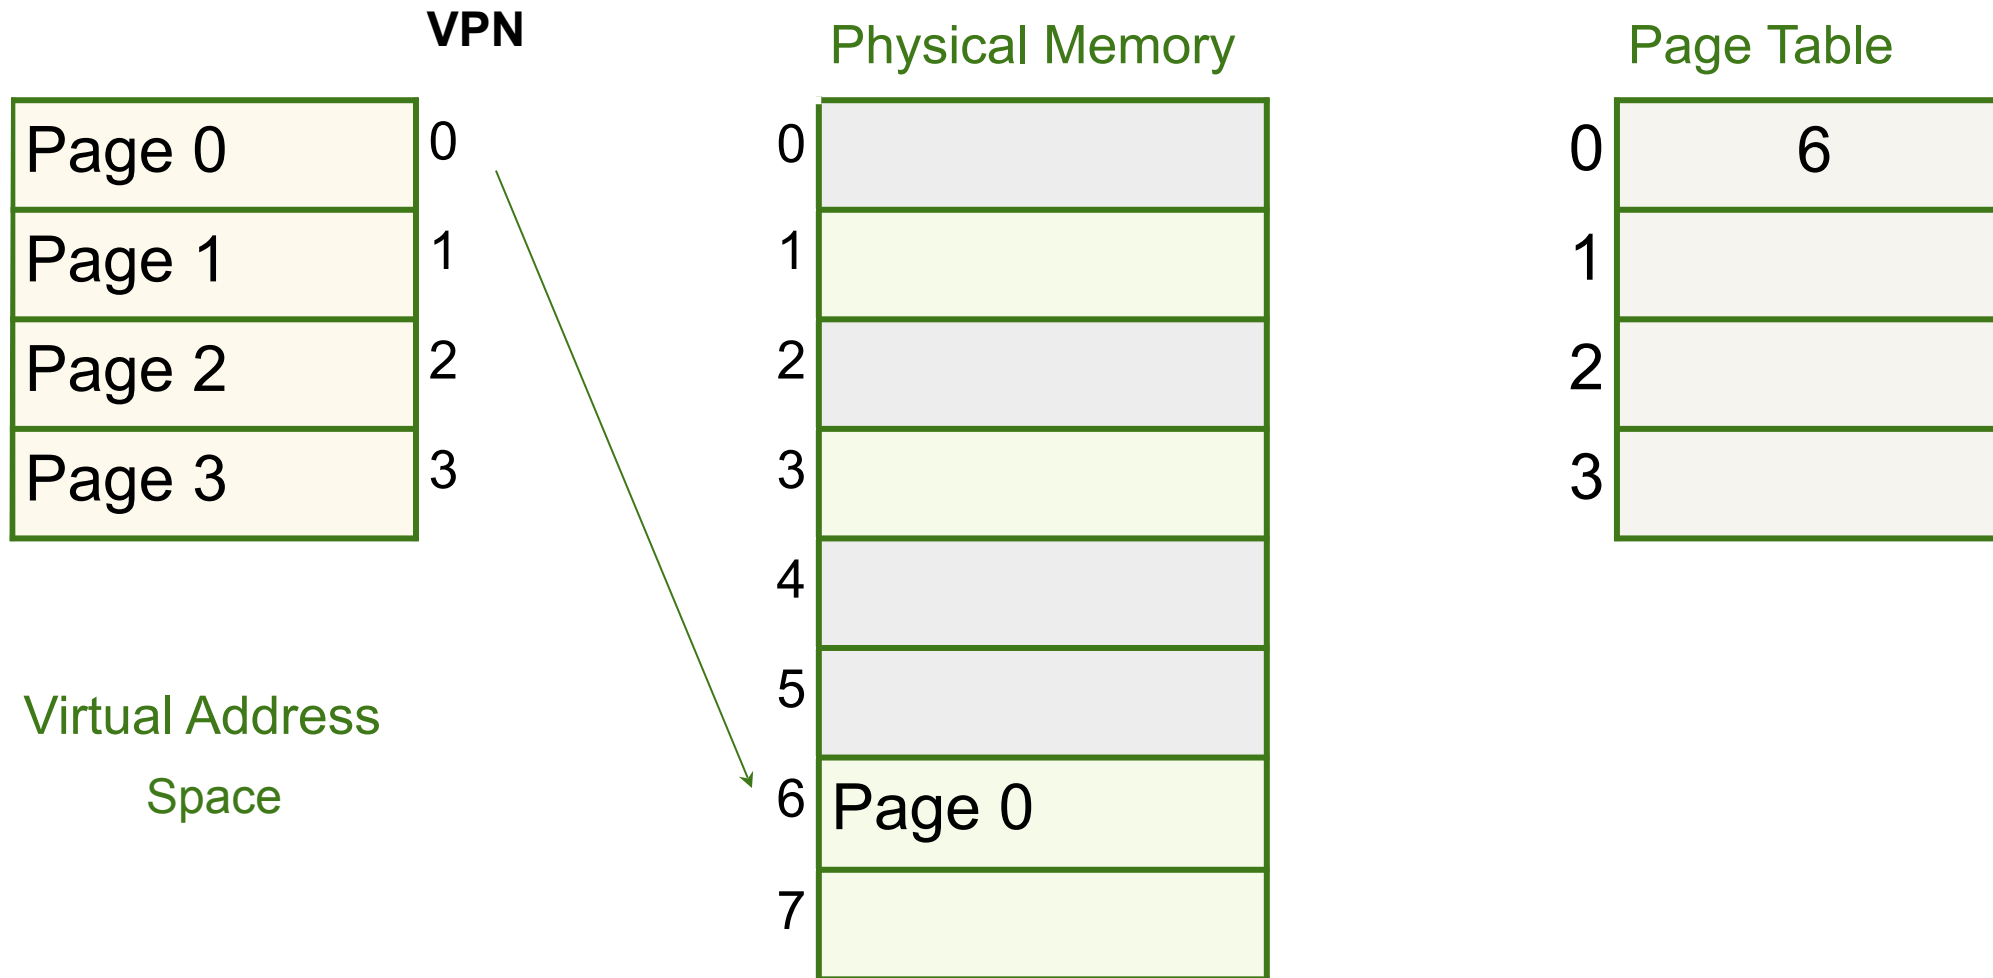
- ❑ This command will open a debugger session for the process with the specified PID. You can then use the `dt nt!_MM_PAGE_TABLE` command to dump the page table for the process.

windbg -i <pid>

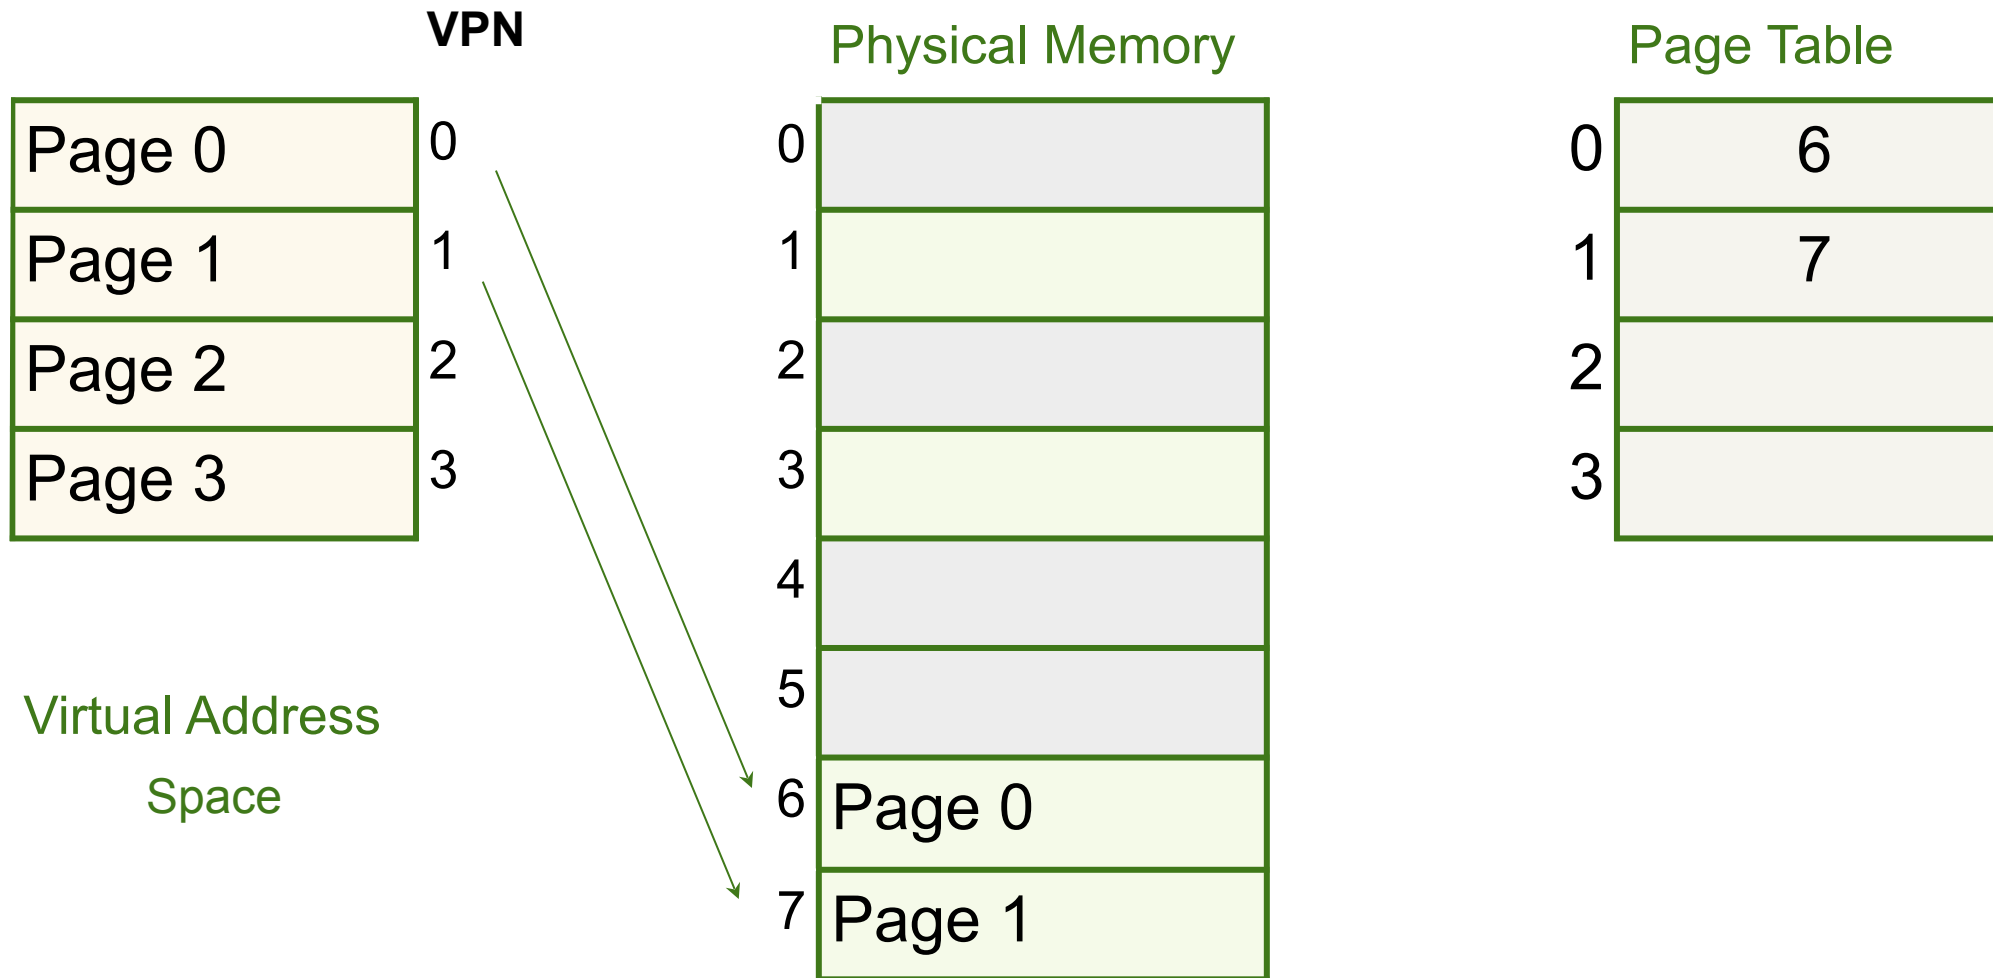
Paging : Memory Allocation



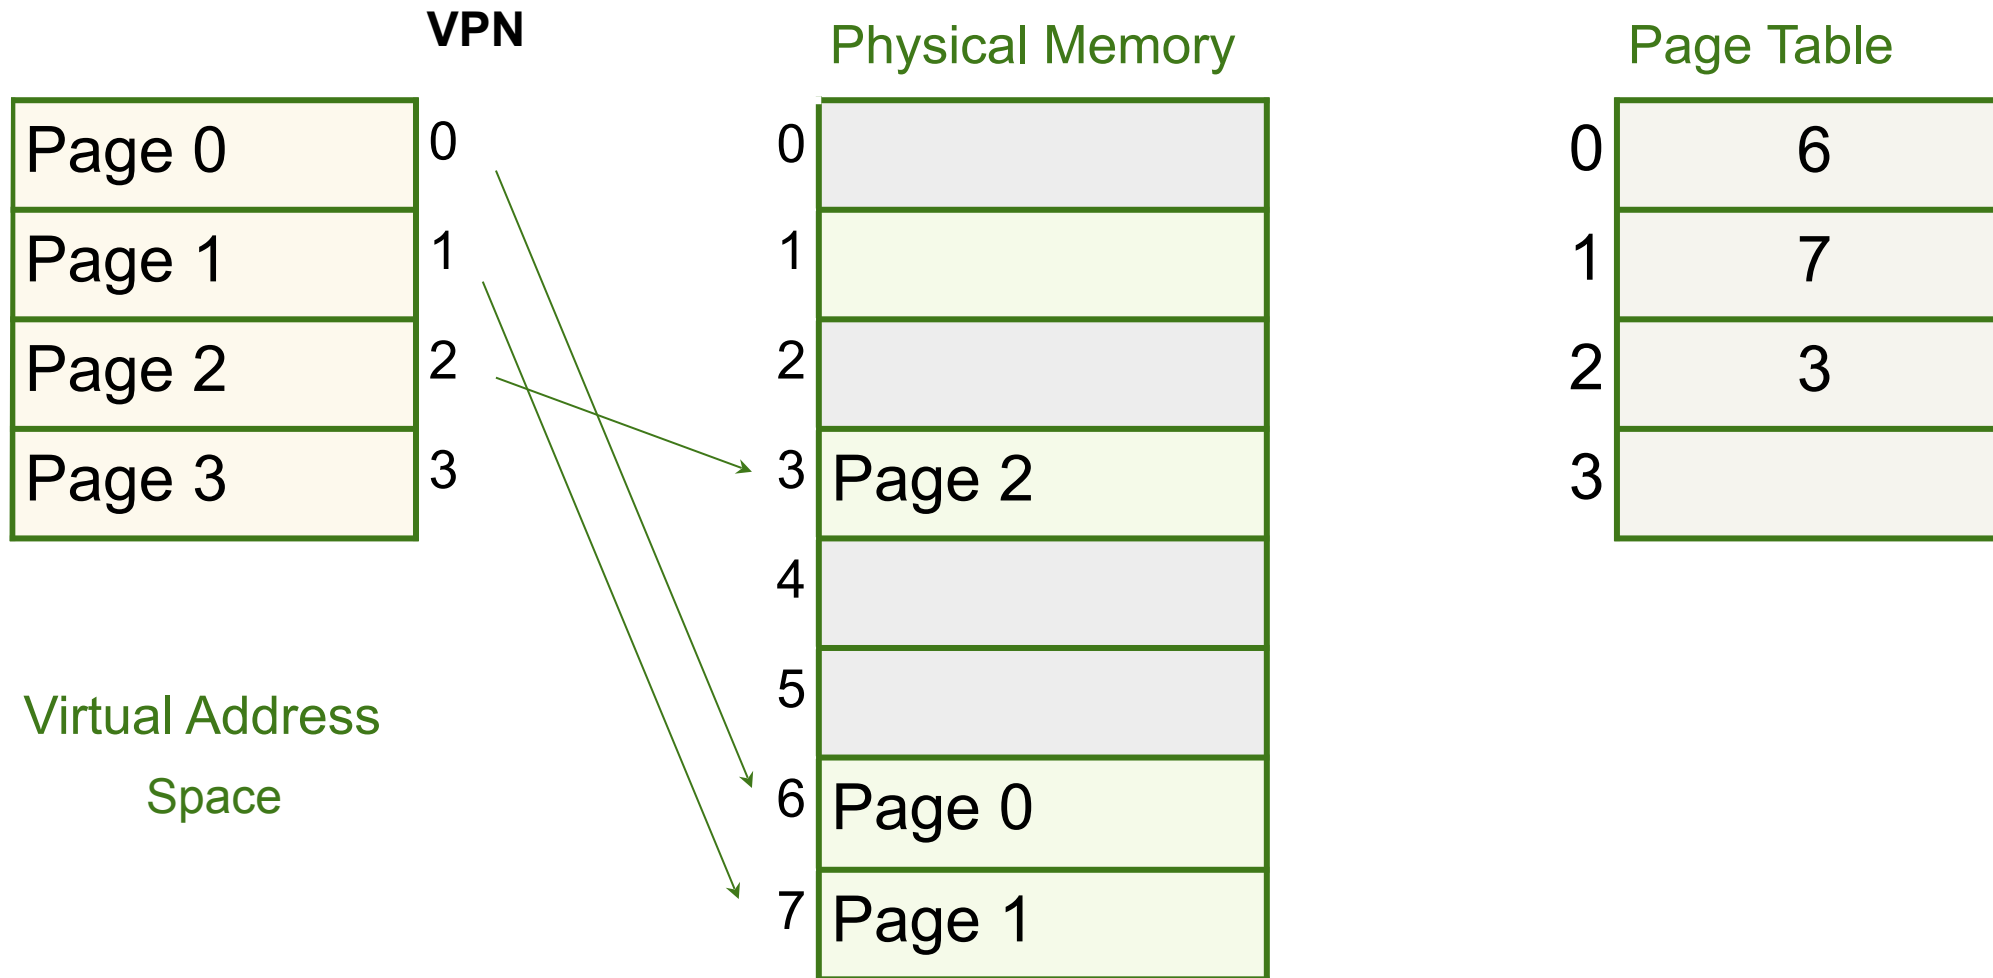
Paging : Memory Allocation




Paging : Memory Allocation



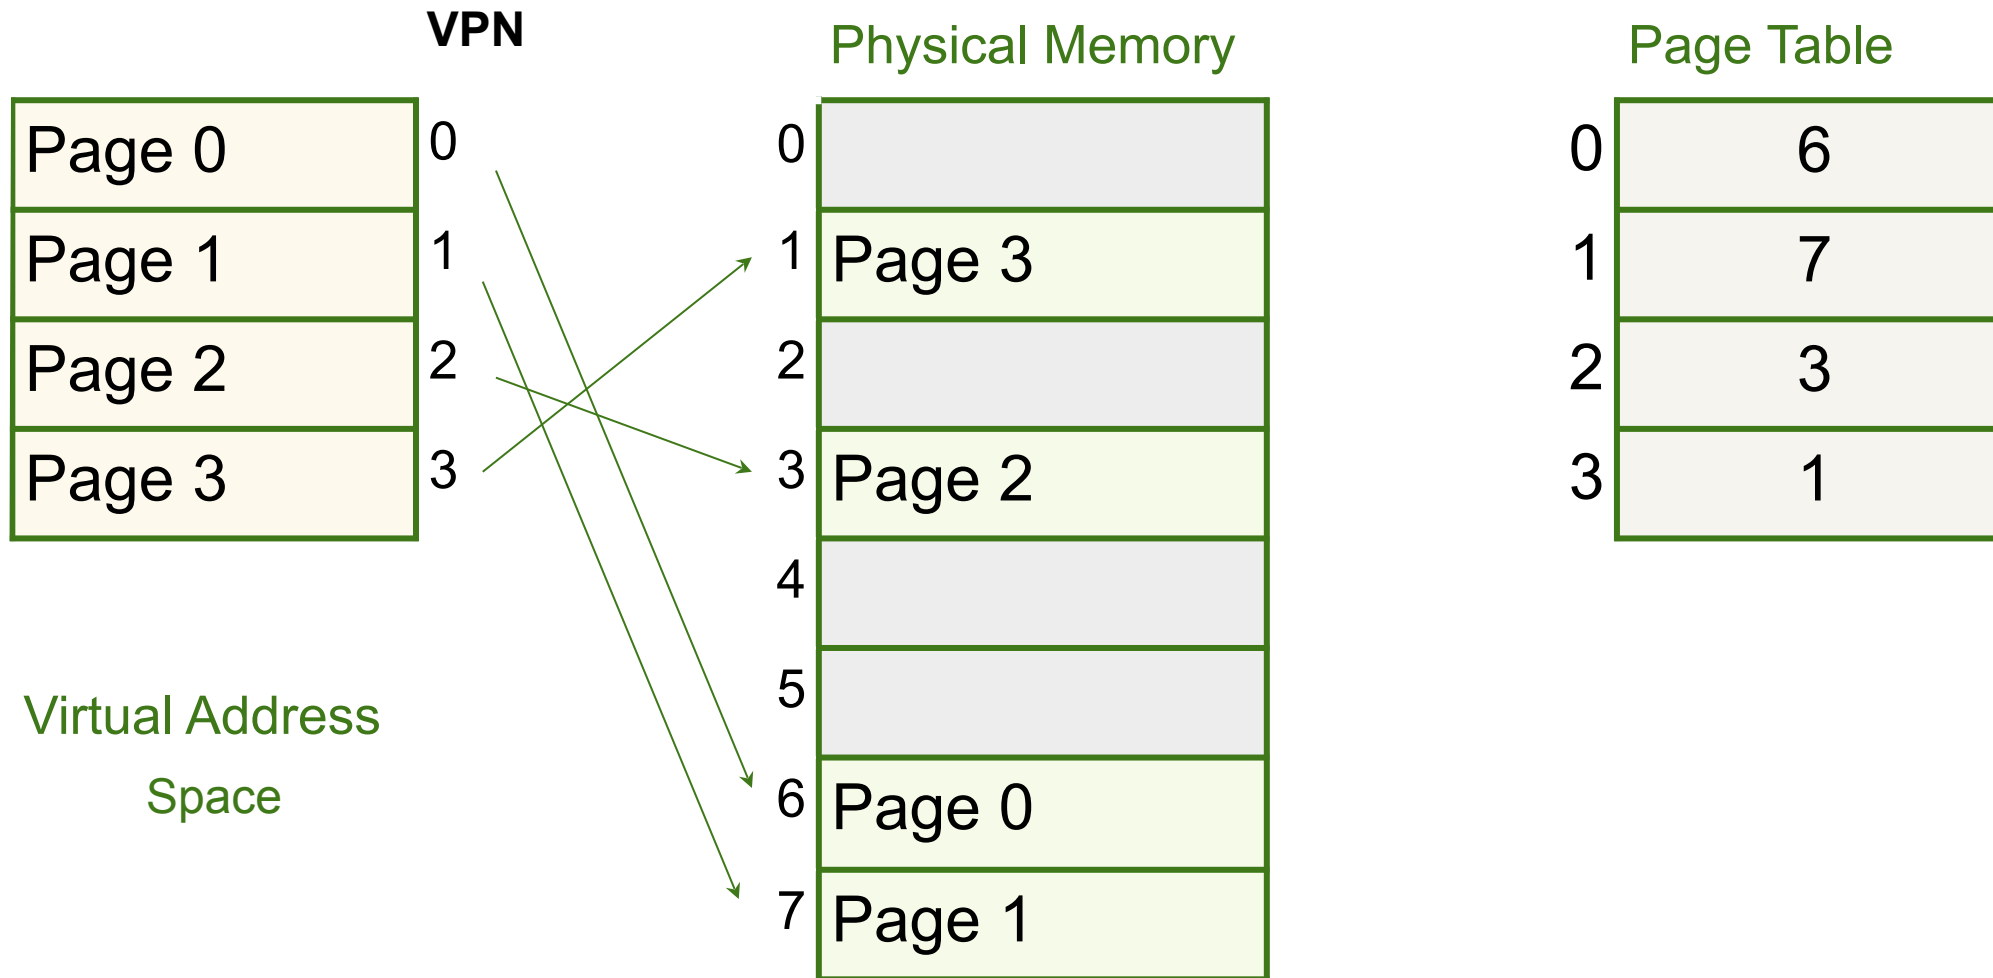
Paging : Memory Allocation

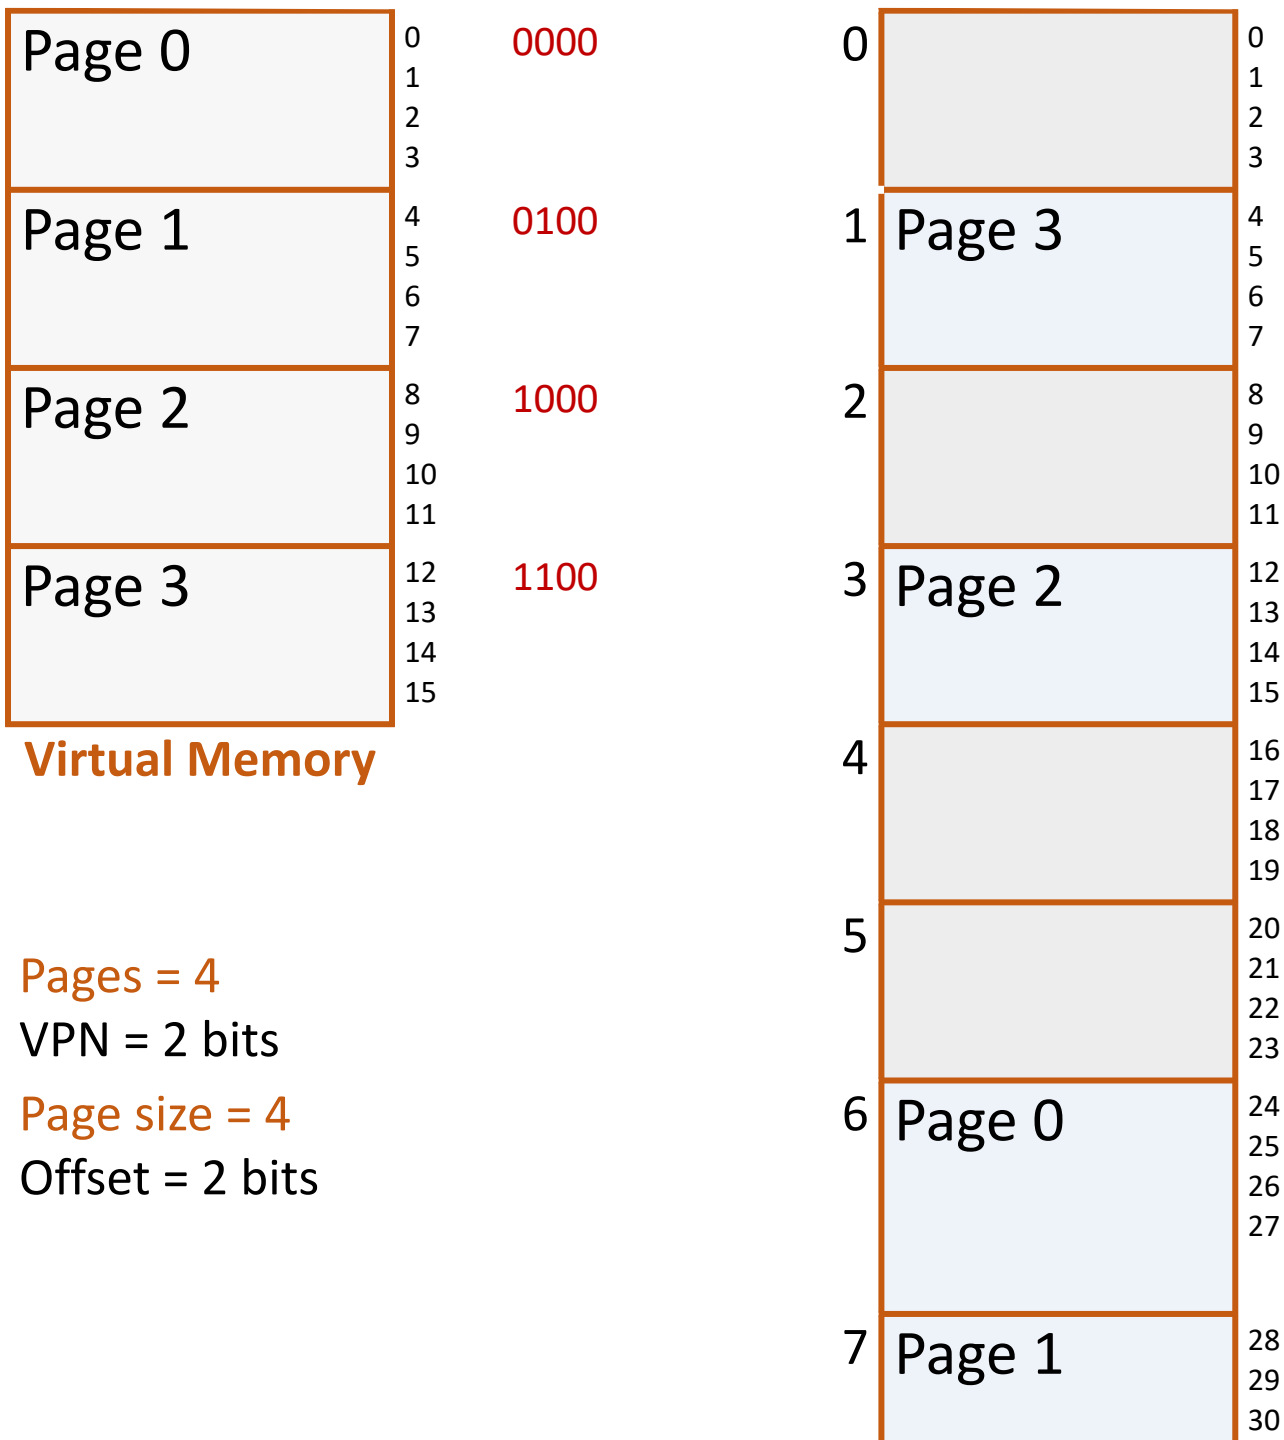


Paging : Memory Allocation



Paging : Memory Allocation





0	6
1	7
2	3
3	1

Page Table

Virtual address = 13 = 1101₂

VPN = 11₂ = 3

Offset = 01₂


No. of Physical Pages = 8

Physical page no = 3 bits

Physical page = 1 = 001₂

Physical address = 00101₂
= 5

Physical Memory



0	6
1	7
2	3
3	1

Page Table

Virtual address = 7 = 0111₂

VPN = 01₂ = 1

Offset = 11₂

No. of Physical Pages = 8

Physical page no = 3 bits

Physical page = 7 = 111₂

Physical address = 111111₂
= 31

Physical Memory

Page Table Entry

- ❑ Page Table Entry consists of
 - Valid bit: indicates if PTE is valid
 - Physical page number
 - Protection bits: read, write, or execute
 - Present bit: page present in memory or swapped to disk
 - Dirty bit: if page contents have changed
 - Reference bits: used in page replacement policy
-

Internal Fragmentation: Calculation

- ❑ Calculating internal fragmentation
 - Page size = 2,048 bytes
 - Process size = 72,766 bytes = 35 pages + 1,086 bytes
 - Internal fragmentation = 2,048 - 1,086 = 962 bytes
 - ❑ Page size = 4KB, Process size = 23MB
 - ❑ Page size = 1024 bytes, Process size = 35,428 bytes
 - ❑ Page size = 4096 bytes, Process size = 4,27,854 bytes
-

Page Table Size Computation

- ❑ Virtual address = 22 bits
 - ❑ Page size = 4 KB
 - ❑ Physical Memory Size = 16 MB
 - ❑ Size of page table?
-

Page Table Size Computation

- ❑ Virtual address = 22 bits
 - ❑ Page size = 4 KB = 2^{10} = 2^{12}
 - ❑ Page offset = 12 bits
 - ❑ No. of bits for VPN = $22 - 12 = 10$ bits
 - ❑ Number of virtual pages = 2^{10}
 - ❑ Number of entries in page table = 2^{10}
 - ❑ Physical Memory Size = 16 MB = $2^4 \times 2^{20} = 2^{24}$
 - ❑ Bits for physical page number = $24 - 10 = 14$ bits
 - ❑ PTE requires 2 bytes
 - ❑ Size of page table = 2×2^{10} bytes = 2KB
-

Page Table Size Computation

- ❑ Virtual address space = 950 MB
 - ❑ Page size = 16 KB
 - ❑ Physical Memory Size = 4 GB
 - ❑ Size of page table?
-

Page Table Size Computation

- ❑ Virtual address = 950 MB = requires 30 bits
 - ❑ Page size = 16 KB = 2^{14} (requires 14 bits)
 - ❑ Physical Memory Size = 4 GB = 2^{32} B
 - ❑ Page offset = 14 bits
 - ❑ No. of bits for VPN = $30 - 14 = 16$ bits
 - ❑ Number of virtual pages = 2^{16}
 - ❑ Number of entries in page table = 2^{16}
 - ❑ Bits for physical page number = $32 - 14 = 18$ bits
 - ❑ PTE requires 3 bytes
 - ❑ Size of page table = 3×2^{16} bytes = 192 KB
-

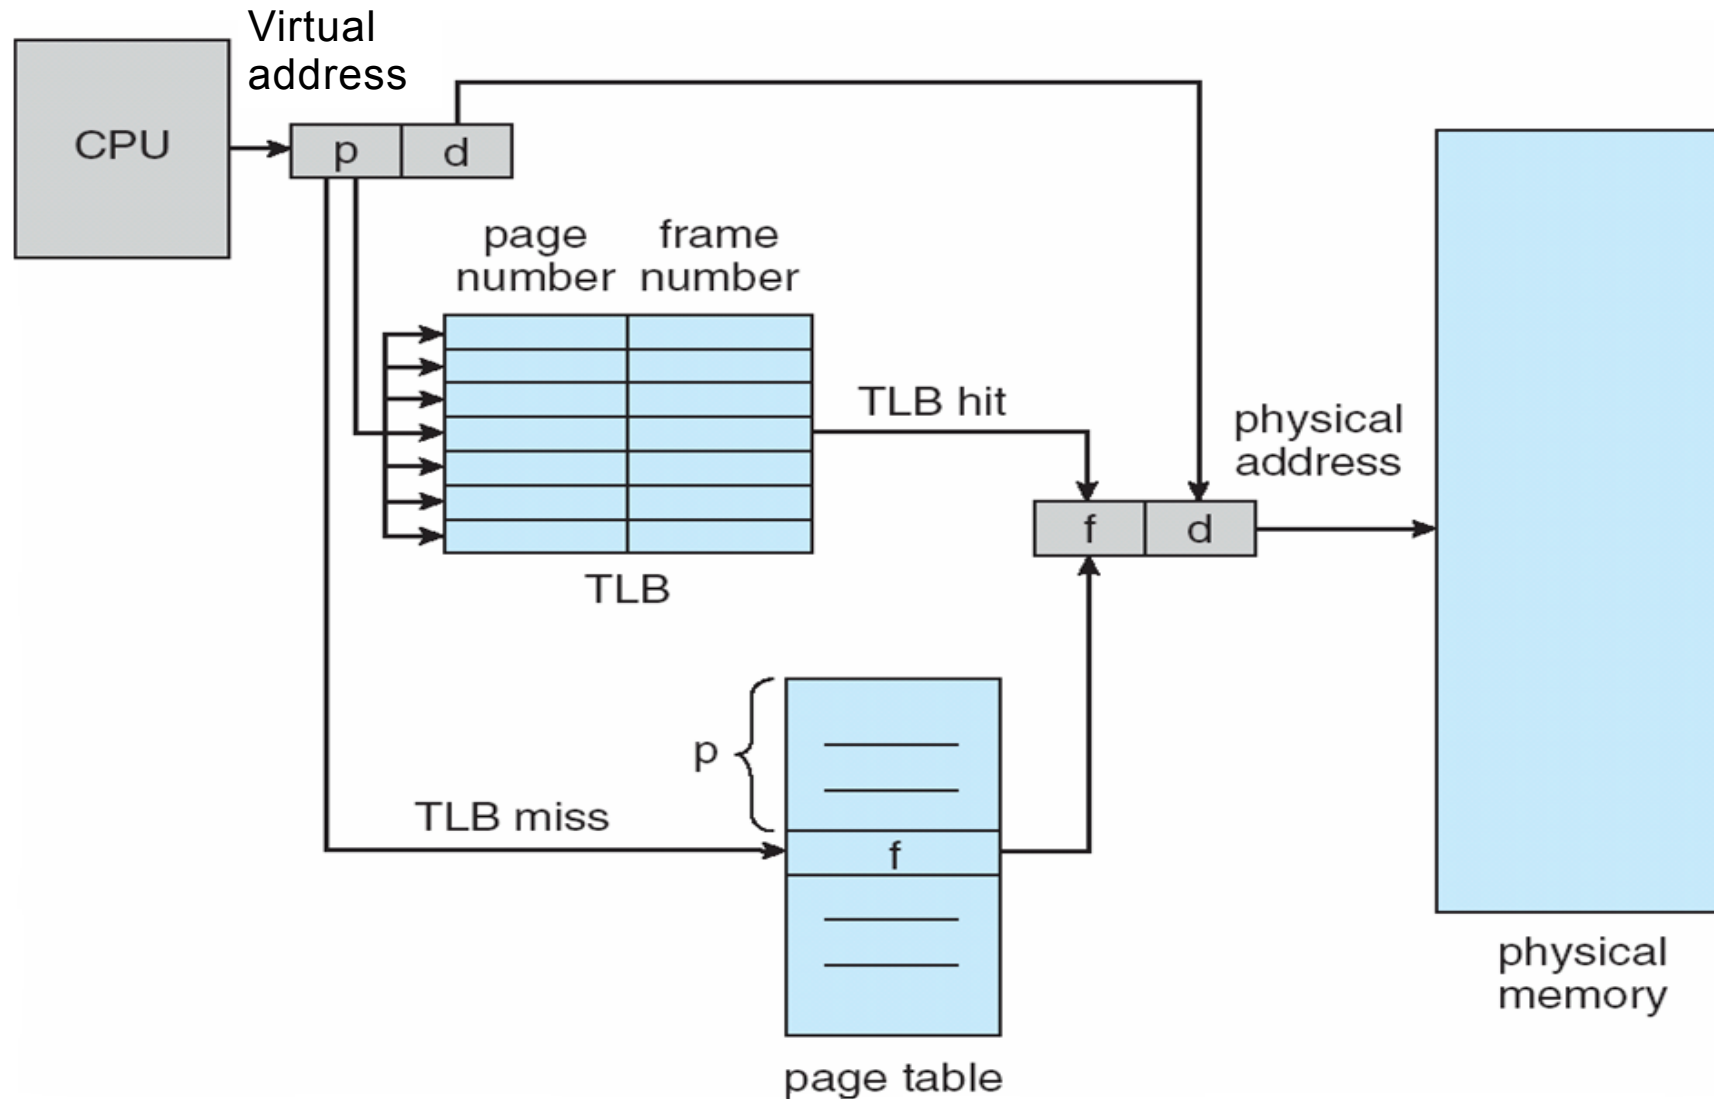
Paging Numericals

- ❑ For an embedded small computer supporting upto 1k bytes physical memory, where its virtual address is 12 bits. Suppose that the size of page/frame is 128 bytes. Use a figure to illustrate the physical and virtual address that shows the number of bits for offset and frame/page number for one level page table.
 - ❑ For an embedded small computer supporting upto 1k bytes physical memory, where its virtual address is 12 bits. Suppose that the size of page/frame is 128 bytes. Use a figure to illustrate the physical and virtual address that shows the number of bits for offset and frame/page number for one level page table.
 - What is the number of virtual pages for each process?
 - How many physical frames in total?
 - How many entries in page table for each process?
-

Page Table Implementation

- ❑ Page table is kept in main memory
 - Page-table base register (PTBR) : starting address
 - Page-table length register (PTLR) : size of page table
 - ❑ Two memory accesses required for every data/instruction
 - First access page table and
 - Second access for the data / instruction
 - ❑ Faster access: special fast-lookup hardware cache called translation look-aside buffers (TLBs) used for page table
 - TLB: associative memory
 - CPU accesses PTE from TLB; on failure page table is accessed
-

Paging: Address Translation with TLB




TLB: Effective Page Access Time

- ❑ Hit ratio = α (percentage of times a valid PTE is found in TLB)
 - ❑ Consider $\alpha = 80\%$,
 - ❑ Access Time $t_A = 20$ ns for TLB,
 - ❑ Access Time $t_M = 100$ ns for memory
 - ❑ Effective Access Time (EAT)
 - $= (t_A + t_M) \alpha + (t_A + 2t_M)(1 - \alpha)$
 - $= 0.80 \times 120 + 0.20 \times 220$
 - $= 96 + 44 = 140\text{ns}$
-

TLB: Effective Page Access Time

- ❑ Hit ratio = α (percentage of times a valid PTE is found in TLB)
 - ❑ Consider $\alpha = 99\%$,
 - ❑ Access Time $t_A = 20$ ns for TLB,
 - ❑ Access Time $t_M = 100$ ns for memory
 - ❑ Effective Access Time (EAT)
 - $= (t_A + t_M) \alpha + (t_A + 2t_M)(1 - \alpha)$
 - $= 0.99 \times 120 + 0.01 \times 220$
 - $= 118.8 + 2.2 = 121$ ns
-

Page Fault



Page Table Size

- ❑ Internal fragmentation (average) = half a page size
 - ❑ Small page size
 - Page table size shall increase
 - Page faults may increase as pages need to be loaded into memory
 - Internal fragmentation shall decrease
 - ❑ Large page size
 - Small page table
 - Memory wasted in internal fragmentation
 - Page fault shall reduce
 - More efficient to transfer to/from disk
-

Thank you.
